

Embedded & Robotics Interview Crash Course


Ordered from fundamentals to senior-level judgment: nine parts, each building on the last — and written for a beginner: every technical term is explained in plain words at its first appearance. Every fact carries its why-chain — because interviewers rarely stop at the first answer — and every MCQ explains why the wrong options feel right.

PART I • BITS & BYTES

Number systems and bit surgery — the alphabet everything else is written in.

1 Hex ↔ Binary Conversion

0x3A



1 hex digit = exactly 4 bits ($16 = 2^4$),
so digits map independently — no carrying,
unlike decimal. That is the whole trick.

Each hex digit converts independently to 4 bits — memorize 16 nibbles, convert anything.

Why? Why does embedded live in hex? First be clear: **0x3A and 58 are the same number** — 0x3A means $3 \times 16 + 10 = 58$; hex and decimal are just two ways of writing one value. The difference is transparency: each hex digit is one nibble, so 0x3A reads directly as 0011 1010 — you see bits 5, 4, 3 and 1 are set at a glance. The decimal form 58 hides that structure completely: to find which bits are set you would have to repeatedly divide by 2. Registers — the chip's fixed-address mailboxes (§16), whose individual bits each mean one on/off thing ("bit fields") — are why datasheets, memory maps and debuggers all speak the transparent base.

→ **Why that?** Why base-16 and not decimal for this? $16 = 2^4$, so digit boundaries align perfectly with bit boundaries. Decimal digits (10 is not a power of 2) smear across bits — no clean mapping.

60-second drill: 0x3A → 0011 1010 · 0xF2 → 1111 0010 · 0xDE → 1101 1110 · reverse: 1010 0101 → 0xA5.

Example — why this matters at the bench: a UART status register reads back as 0x3A. Split the digits: 3 → 0011, A → 1010, so bits {5, 4, 3, 1} are set — and the datasheet's bit map turns that straight into "TX-empty, TX-complete, RX-full, overrun" without a single calculation. The same register printed as decimal 58 tells you nothing until you convert. That gap — read the bits vs compute the bits — is the entire reason every debugger defaults to hex.

Q: 0xF2 in binary?

- A. 1111 0010 ✓ *F = 1111, 2 = 0010. Two independent nibbles, done.*
- B. 1010 1010 ✗ *That is 0xAA — a favorite test pattern, so it looks familiar. Familiar / correct.*
- C. 0000 1111 ✗ *That is 0x0F. Tempting if you only converted the "F" and lost track of digit position.*
- D. 1111 1111 ✗ *That is 0xFF. The "F is all-ones" reflex applied to both digits — check every nibble separately.*

Interview tip: Bonus they ask next: "What does $x \ \& \ (x - 1)$ do?" — clears the lowest set bit; zero result means x was a power of two.

2 Data Types: uint8_t vs uint32_t

uint8_t (1 byte) value 0x3A = 58



max 255 — one CAN byte, ADC8, flags

uint32_t (4 bytes)



max 4 294 967 295 — timestamps, encoder counts, addresses. Shaded: same 8-bit payload wastes 24 bits here.

Same payload, different footprint: pick the smallest type that can never overflow.

Why? Why care about 3 saved bytes? Multiply by a 10,000-element sensor buffer, a CAN frame (§23's vehicle-bus parcel) that fits 8 bytes total, or an EEPROM (a small on-chip memory that survives power-off) with 1 kB — width discipline is what makes small systems possible.

→ **Why that?** Why not always shrink to uint8_t? On a 32-bit CPU, 8-bit arithmetic can be *slower* (masking to emulate wrap-around), and C integer promotion silently widens operands anyway. Rule: small types for **storage and protocols**, natural width (int / uint32_t) for **loop counters and math**.

→ **Why that?** Why do overflows here become interview war stories? uint8_t count = 255; count++; wraps to 0 with no warning — the classic "counter reset itself" field bug.

Q: A wheel encoder produces up to 2 million counts per revolution accumulation. Best type for the running count?

A. uint8_t **X** *Smallest = most efficient, right? Wraps at 255 — the count would reset ~7,800 times per rev.*

B. uint16_t **X** *Feels like a sensible middle ground, but 65,535 max still overflows 30× per revolution.*

C. uint32_t **✓** *Headroom to 4.29 billion, native width on Cortex-M — free performance, no overflow in practical lifetime.*

D. float **X** *Tempting because "big numbers" — but floats lose integer exactness above 2²⁴ (16.7 M): counts start silently skipping.*

3 The Bit-Manipulation Toolkit

The bit-manipulation family — toggle was the test question ("no if/ternary allowed"), but learn all four as one unit:

```
v |= (1UL << pos);      /* SET bit      */
v &= ~(1UL << pos);     /* CLEAR bit   */
v ^= (1UL << pos);      /* TOGGLE bit  -- XOR flips, no branch needed */
bit = (v >> pos) & 1u;  /* READ bit    */

uint32_t toggle_bit(uint32_t value, uint8_t pos) {
    return value ^ (1UL << pos);
}

/* worked — each line starts fresh from v = 0x3A = 0011 1010 :
v |= (1UL << 6);      /* mask 0100 0000 → v = 0x7A = 0111 1010: bit 6 set */
v &= ~(1UL << 3);     /* mask 1111 0111 → v = 0x32 = 0011 0010: bit 3 gone */
v ^= (1UL << 5);      /* mask 0010 0000 → v = 0x1A = 0001 1010: bit 5 flip */
v ^= (1UL << 5);      /* toggle again   → v = 0x3A: XOR twice = restored */
bit = (v >> 5) & 1u; /* 0x3A >> 5 = 0000 0001 → bit = 1: bit 5 was set */
/* and the function: toggle_bit(0x3A, 0) == 0x3B — bit 0 was 0, now 1 */
```

On real hardware — the same four lines are how every driver touches a peripheral (a built-in helper circuit — §16; GPIOA is the register block controlling the chip's port-A pins):

```

GPIOA->ODR |= (1u << 5);      /* LED on PA5: read-modify-write */
GPIOA->ODR ^= (1u << 5);      /* toggle it */
GPIOA->BSRR = (1u << 5);      /* set via BSRR — a single WRITE: atomic */
GPIOA->BSRR = (1u << (5 + 16)); /* upper half of BSRR clears the pin */
/* why BSRR exists: |= is load-modify-store, and an ISR between the load and */
/* the store loses a bit (§39's i++ problem). One write cannot be interrupted. */

```

Why? Why does XOR toggle? $x \text{ XOR } 0 = x$, and $x \text{ XOR } 1 = \text{NOT } x$ — the mask's 1 flips exactly one bit, its 0s pass everything else through. The "no if" constraint is the hint: they want proof you think in masks, not branches.

→ **Why that?** Why 1UL and not 1? Plain 1 is a signed int; on a 32-bit int, $1 << 31$ is undefined behavior (shifting into the sign bit). 1UL makes the mask unsigned before shifting — a one-character answer that separates candidates who have shipped from those who haven't.

PART II • C, THE INTERVIEW LANGUAGE

Pointers, qualifiers and memory — where most screening rounds are actually decided.


4 Pointers & Pass-by-Reference — the Classics

The pass-by-reference trap. Fill in the missing line so the caller's variable actually changes:

```

void increment_by(int *value, int amount) {
    /* ??? */
}
int x = 10;
increment_by(&x, 5); /* x must become 15 */

```



The trap in one picture: arithmetic on p moves the arrow; *p follows it.

Q: Which line is correct?

- A. `value = value + amount;` X *Pointer arithmetic on the local copy of the pointer: it now points 5 ints (20 bytes) away, x untouched. Compiles cleanly — the worst kind of wrong.*
- B. `value += amount;` X *Identical bug in compact clothing — the += makes it look like the intended increment, which is exactly why it's offered.*
- C. `*value = *value + amount;` ✓ *Dereference to reach the caller's int, read it, add, write it back. (*value) += amount; is equivalent.*
- D. `*value = value + amount;` X *Assigns an address+offset into an int — type confusion the compiler warns about. If you can't say why each side's type differs, revisit §5.*

Interview tip: The interviewer's follow-up: "what does the caller pass?" — `&x`, an address. Parameters are always copies; a copy of an address still reaches the original object. That sentence is the whole concept.


Pointer arithmetic scales by the pointed-to type — `p+1` means "the next object", never "the next byte":

```

int32_t v[3];           /* say v lives at 0x2000_0000 */
int32_t *pi = v;
pi + 1;                 /* = 0x2000_0004 - stepped sizeof(int32_t) = 4 */
((uint8_t *)v) + 1;     /* = 0x2000_0001 - same start, 1-byte stride */

&v[2] - &v[0];          /* == 2 - subtraction yields ELEMENT counts */
/* (ptrdiff_t), defined only within one array */

```



same memory, same +1 in the source — the TYPE decides the stride

One +1 in the source, two different jumps in memory — the declared type is the unit.

Why? Why does C scale instead of counting bytes? So that indexing can be arithmetic: $a[i]$ is defined as $*(a + i)$, which only lands on element i if the $+$ moves in whole elements. It is also exactly why §4's trap moved 20 bytes: $\text{value} + 5$ on an int pointer is $5 \times \text{sizeof}(\text{int})$.

Arrays and pointers — intimate, not identical. In almost every expression an array name decays to a pointer to its first element — which is why $a[i]$, $*(a+i)$ and even $i[a]$ are the same thing. The identity breaks in three places: $\text{sizeof}(a)$ is the whole array, $\&a$ is a pointer to the array (stride = whole array), and you cannot assign to an array. The cruelest decay happens at function boundaries — a parameter declared `int a[10]` is silently rewritten to `int *a`, and the length is gone: pass it explicitly, always. Two declarations to read right-to-left (§5's skill again):

```

int *p[5];           /* array of 5 pointers-to-int */
int (*p)[5];         /* pointer to one array-of-5-ints - p+1 jumps 20 B */

```

Q: `void f(int a[10]) { return sizeof(a); }` — what does it return on a 32-bit MCU?

- A. 40 — ten ints of four bytes **✗** What the declaration looks like it promises. But array parameters decay: the `[10]` is documentation the compiler discards.
- B. 10 — the element count **✗** Doubly wrong: `sizeof` never counts elements, and the array-ness is already gone.
- C. 4 — `sizeof(int *)` **✓** Inside `f`, `a` IS a pointer. This is why every array-taking API carries a separate length parameter — and most compilers warn about this exact `sizeof`.
- D. It does not compile **✗** Legal (if suspicious) C — which makes it a better interview question than a compile error would.

The two-strings trap — same characters, opposite memory (§7's map decides):

```

char *s = "hello";    /* pointer into the string LITERAL in .rodata/flash */
char t[] = "hello";    /* 6-byte writable COPY on the stack */

s[0] = 'H';           /* fault (or silent nothing): writing to flash */
t[0] = 'H';           /* fine — the copy is yours */
/* also: sizeof(s) == 4 (a pointer), sizeof(t) == 6 (the array + NUL) */

```

Double pointers (`T **`) — §4's rule applied one level up: to change the caller's *value* you pass its address; to change the caller's *pointer* you pass the pointer's address. The canonical need is a linked list's head — and it is why `main` receives `char **argv`:

```
void push_front(node_t **head, node_t *n) {
    n->next = *head;    /* new node aims at the old first node    */
    *head = n;          /* and the CALLER's head now aims at the new */
}
/* with a plain node_t *head parameter, only a local copy would move */
```

Reverse an array in place (no second array) — the two-pointer swap:

```
void reverse_array(int32_t *a, uint32_t n) {
    for (uint32_t i = 0, j = n - 1; i < j; i++, j--) {
        int32_t tmp = a[i];
        a[i] = a[j];
        a[j] = tmp;
    }
}
/* why i < j and not i <= j? meeting in the middle: equal index = nothing to swap; */
/* crossing = re-swapping back. Also: what if n == 0? n-1 wraps to 4 billion (uint!) — */
/* guard with if (n < 2) return; Saying this unprompted is the senior signal. */
```

Reverse a singly linked list — the pointer question coding rounds standardize on:

```
node_t *reverse(node_t *head) {
    node_t *prev = NULL;
    while (head) {
        node_t *next = head->next; /* save the rest BEFORE cutting it off */
        head->next = prev;         /* flip one arrow */
        prev = head;              /* both walkers advance */
        head = next;
    }
    return prev;                 /* old tail is the new head */
}
```

Why? Why three pointers? Flipping `head->next` destroys the only road forward — `next` is the bookmark that saves it first. One pass, $O(1)$ extra space; the follow-up is nearly guaranteed: “detect a cycle” → Floyd’s slow/fast walkers — if they ever meet, the list loops.

Sharp edges to name before they cut you: comparing or subtracting pointers from *different* arrays is undefined behavior, not just rude. Casting a `uint8_t` buffer to `uint32_t*` to read a protocol field risks an unaligned access (\$8’s padding story) *and* violates strict aliasing — the sanctioned move is `memcpy` into a local, which optimizers turn into the same load anyway. And `void*` has no arithmetic in standard C — GCC treats it as byte-sized as an extension, which is exactly the kind of portability landmine a “why?” question is fishing for.

5 const and Pointers — Read Right to Left

The technique first. Start at the variable name and read **right to left**, with one binding rule: **const modifies whatever sits immediately to its left** — and only if nothing is there does it take the thing on its right. That single rule makes the whole table mechanical instead of memorized.

Walk `const int *const p` right to left: `p` ... is a const pointer ... (*) ... to an `int` that is const. Two consequences fall out immediately. First, `const int *p` $\equiv$ `int const *p` — in both, the `const` sits on the *int* side of the `*`, so both freeze the **data**; only a `const` to the *right* of the `*` (touching `p`) freezes the **pointer**. Second, that is exactly why `p++` is legal in the table’s first row: the pointer variable itself was never declared `const` — only what it points at.

Declaration	Read it as	What is frozen
<code>const int *p</code>	<code>p</code> is a pointer to a const int	The data : <code>*p = 5</code> → compile error; <code>p++</code> is fine
<code>int *const p</code>	<code>p</code> is a const pointer to int	The address : <code>p = &x</code> → error; <code>*p = 5</code> is fine
<code>const int *const p</code>	const pointer to const int	Both. Typical for memory-mapped ROM tables

Why? Why does embedded care so much? Register definitions: `volatile uint32_t *const UART_DR` — the peripheral address never moves (const pointer) but its contents change under you (hence not const data, hence volatile).

Now watch the compiler enforce each row — the same three declarations, exercised:

```
int a = 1, b = 2;

const int *p = &a;          /* row 1: pointer TO const data          */
*p = 5;                     /* ERROR: assignment of read-only location */
p = &b;                     /* fine — the pointer itself may move      */

int *const q = &a;          /* row 2: CONST pointer to writable data   */
*q = 5;                     /* fine — a is now 5                      */
q = &b;                     /* ERROR: assignment of read-only variable */

const int *const r = &a;    /* row 3: both frozen                      */
/* only reading *r is legal — everything else is a compile error */
```

The register idiom, spelled out — the line every HAL is built from, now readable right-to-left:

```
volatile uint32_t *const UART_DR = (uint32_t *)0x40011004;
/* read it: UART_DR is a CONST pointer (the peripheral never moves, */
/* so the compiler may hard-code the address into every access)      */
/* to a VOLATILE uint32_t (hardware rewrites it behind your back, §6) */

*UART_DR = 0x41;            /* write → transmit 'A'                    */
c = *UART_DR;              /* read → receive — re-read every time    */
UART_DR = &elsewhere;      /* ERROR — peripherals do not wander      */
```

And the API meaning: `void send(const uint8_t *buf, size_t n)` is a promise — “I will read your buffer, never write it” — which lets callers hand over flash-resident data safely and turns any accidental write inside `send()` into a compile error. One sentence covers both interview angles: **const on a parameter is a contract; const on a file-scope object is also a placement hint** — it is what lets the linker put data in `.rodata`/flash (§7).

Q: A 256-entry CRC lookup table must never be modified and should live in flash, not RAM. Best declaration?

- A. `uint8_t table[256]` ❌ Writable — so it lands in `.data`: a RAM copy initialized at boot (§12), wasting exactly the 256 bytes you meant to save.
- B. `const uint8_t table[256]` ✅ A const file-scope object goes to `.rodata` — flash on an MCU — and every write attempt is a compile error. Placement and protection from one keyword.
- C. `uint8_t *const table` ❌ Freezes the wrong thing: the POINTER is const, the bytes it aims at stay writable — and stay in RAM. Right-to-left reading catches this instantly.
- D. `volatile const uint8_t table[256]` ❌ Legal but self-defeating: volatile forbids caching reads of data that, being const, can never change — you pay for freshness nothing can supply.

6 volatile

Tells the compiler: *this value can change outside normal program flow* — never cache it in a register, never optimize the access away, re-read it every time.

Why? Why is it needed at all? The optimizer (the compiler stage that rewrites your code to run faster) sees `while(flag == 0);`, proves no code in the loop writes `flag`, hoists the read, and spins forever on a stale copy. It cannot see the ISR (the short function hardware runs the instant an event fires — §12) or the hardware register that writes it.

→ **Why that?** Why the three canonical uses? (1) memory-mapped peripheral registers, (2) variables shared with ISRs, (3) memory touched by DMA (§16’s no-CPU data courier) or by other cores — all writers invisible to the compiler.

→ **Why that?** Why is volatile **not** enough for thread safety? It forbids caching, but guarantees neither atomicity nor ordering. A 32-bit write on an 8-bit MCU is 4 separate stores an ISR can interrupt halfway. You still need critical sections or atomics — saying this unprompted is a strong-hire signal.

The terms doing the heavy lifting — unpacked, with examples

“Cache it in a register” / “hoist the read”. Memory is slow and registers are fast, so the optimizer loads a variable **once** into a CPU register and reuses that copy — like copying a phone number onto a sticky note: fast, but if the phonebook changes, the note never hears about it — and it “hoists” loads out of loops whenever nothing *inside the loop* writes the variable. This is the silent transformation:

```
/* you wrote:                what the optimizer emits (no volatile): */
while (flag == 0) {          r0 = flag;           /* ONE load, hoisted */
                             while (r0 == 0) {} /* spins on the */
                             /* REGISTER copy - */
                             /* forever      */
```

Atomicity = an operation that completes as one indivisible step — nothing can ever observe it half-done. Everyday version: changing a 4-dial lock code one dial at a time while someone peeks halfway — they read a code that never existed. Most C statements are *not* atomic:

```
uint32_t count = 0xAAAAAAA; /* on an 8-bit AVR, one C line = 4 stores */
count = 0x11223344;
/* store 0x44 ... store 0x33 ... [ISR fires HERE, reads count] ... 0x22 ... */
/* the ISR saw 0xAAAA3344 — a TORN value that never existed. */
/* Same disease as i++ (load-add-store, §39). volatile does NOT cure it: */
/* it forces the accesses to happen, not to happen indivisibly. */
```

Ordering = the sequence in which memory operations become visible to someone else. The compiler (and, on larger cores, the CPU itself) may reorder accesses that look independent to it — a postal system free to deliver the parcel and the letter announcing it in either order:

```
data = 42; /* you meant: payload first... */
ready = 1; /* ...then the flag. Nothing FORCES that order for plain */
          /* variables — a consumer can see ready==1 with stale data. */
/* volatile preserves order only AMONG volatile accesses; plain data may */
/* still float past it. Real handshakes use barriers or C11 atomics. */
```

Critical section = a stretch with interrupts masked, so a multi-step access completes as a unit — a “do not disturb” sign hung on the door for exactly as long as the delicate job takes. Atomicity bought by paying with interrupt latency (§40), which is why it must stay a few instructions long:

```
uint32_t m = __get_PRIMASK();
__disable_irq(); /* nothing can interleave from here... */
shared += 1;      /* the load-add-store now completes whole */
__set_PRIMASK(m); /* ...to here. Restore, never blindly enable. */
```

Atomics = hardware doing the same job without masking anything: C11 `atomic_fetch_add(&shared;, 1)`, implemented on ARM as an LDREX/STREX pair — the exclusive store *fails* if anyone touched the location in between, and the loop simply retries — like signing a form only if the page is untouched since you read it, else fetch a fresh copy and try again. With these five terms unpacked, the why-chain above stops being vocabulary and becomes mechanics.

The minimal example this whole section is about — one flag between an ISR and main:

```
volatile uint8_t rx_ready;           /* ISR writes, main reads      */

void UART_IRQHandler(void) { data = UART->DR; rx_ready = 1; }

for (;;) {
    while (!rx_ready) { }           /* WITHOUT volatile: the read is    */
    rx_ready = 0;                   /* hoisted and this becomes while(1) */
    handle(data);
}
```


Q: Why declare a variable volatile?

- A. To prevent the compiler optimizing away accesses ✓ *The definition: force a real memory read/write every time.*
- B. To save memory ✗ *Opposite if anything — it blocks register-caching optimizations.*
- C. To make code faster ✗ *Tempting because keywords "tune" code — volatile strictly slows access by forbidding caching.*
- D. To make the variable global ✗ *Confuses storage qualifiers with scope — volatile says nothing about visibility.*

7 Stack vs Heap & the Memory Map

Two kinds of working memory, two personalities: the **stack** is automatic scratch space that grows with every function call and vanishes on return; the **heap** is memory you explicitly rent with `malloc()` and must remember to return with `free()`. The map shows where both live — and why their collision is the classic crash:

RAM (typical MCU memory map)



MCU memory map: stack and heap grow toward each other across the free gap — the collision is your crash.

	Stack	Heap
Allocation	Automatic — move the stack pointer (1 instruction)	malloc/free walk a free-list (hundreds of cycles, variable)
Lifetime	Dies with the function	Yours until freed — and yours to leak
Failure mode	Overflow: silently smashes neighboring memory (or the small per-task stacks an RTOS — Part VI's mini operating system — gives each task)	Fragmentation: free bytes exist but no contiguous block → malloc fails after weeks of uptime
Determinism	Fully deterministic	Non-deterministic timing — hostile to real-time
Embedded verdict	Default choice; size each RTOS task stack deliberately	Many firmware standards allow it only at init, or ban it (MISRA)

Why? Why is fragmentation fatal for devices that run for years? Desktop apps restart daily; a smart meter runs for a decade. Allocation patterns that "work in testing" fragment slowly until one malloc fails in the field — with no user around to reboot it.

Example of the classic overflow — an innocent helper on a 512-byte task stack:

```
void log_line(const char *msg) {
    char buf[600];                /* lives on the CALLER's task stack */
    snprintf(buf, sizeof buf, "[%lu] %s", tick, msg);
} /* 600 > 512: quietly smashes the neighbor — the crash appears later, */
/* somewhere unrelated. Exactly the bug $41's Question 1 is about. */
```

8 The C Questions Testimonials Keep Reporting


Scrape interview write-ups and the same C questions surface again and again. This section is that list, answered.

Struct padding & alignment — the single most-reported sizeof trap:


```
struct S {
    char a; /* 1 byte, then 3 padding bytes */
    int b; /* must start on a 4-byte boundary */
    char c; /* 1 byte, then 3 tail-padding bytes */
}; /* sizeof(struct S) == 12, not 6 */

struct T { /* same members, reordered: */
    int b; /* 4 */
    char a; /* 1 */
    char c; /* 1, then 2 tail bytes */
}; /* sizeof(struct T) == 8 — ordering matters */
```

struct S { char a; int b; char c; } — sizeof = 12



struct T { int b; char a; char c; } — same members, sizeof = 8



red cells are pure alignment tax — int b must start at a multiple of 4, and arrays need the tail rounded up

Same members, two layouts — padding is the alignment tax, ordering is the refund.

Why? Why does the compiler waste bytes? The bus fetches aligned words: a 4-byte int starting at address 1 would need two memory cycles to assemble — or, on many ARM cores, simply hard-faults. Padding buys single-cycle access with memory you usually have; the compiler chose speed and correctness over size.

→ **Why that?** Why does `__attribute__((packed))` exist then? For data that must match an external layout bit-for-bit — a protocol frame, a flash record. The bill: the compiler emits byte-by-byte access code, slower and bigger; take a pointer into a packed struct and hand it to normal code, and you have re-created the unaligned fault. Rule: pack at the wire/storage boundary, copy into natural structs to work.

→ **Why that?** Why is the tail padding there (the 12 vs 10)? Arrays. `S arr[2]` must keep every element's int aligned, so the struct is rounded up to a multiple of its strictest member's alignment.

Q: On a 32-bit MCU, sizeof(struct { char a; int b; char c; }) is...

A. 6 — the sum of the members **✗** *The intuitive answer, and the whole reason this question exists: it assumes memory is byte-addressable soup with no alignment cost.*

B. 8 — the compiler packs a and c together **✗** *Half-learned padding: the compiler may never reorder members (the standard forbids it), so a's padding and the tail both stay.*

C. 12 — 3 bytes after a, 3 after c **✓** *Alignment before b, tail padding for arrays. Follow-up you should volunteer: reordering to int-first shrinks it to 8.*

D. 16 — structs round up to a power of two **✗** *An invented rule stated confidently — rounding is to the largest member's alignment (4 here), not to powers of two.*

The pointer taxonomy — four adjectives, one line each: a **null** pointer points at a defined nowhere (safe to test, fatal to dereference); a **wild** pointer was never initialized (points anywhere); a **dangling** pointer was valid until its object died (freed heap block, returned local's address) — the nastiest, because it often still “works”; a **void*** is a typeless address for generic code — what malloc returns and memcpy accepts — that must be cast back before arithmetic or dereference.

Function pointers — the syntax everyone blanks on, and why embedded cannot avoid it: the vector table you will meet in §12 is an array of function pointers, and every callback/driver-ops pattern is built from them.

```
void (*handler)(int) = my_isr; /* pointer to: function taking int, returning void */
void (*table[8])(void);      /* the vector-table shape: array of function ptrs */
typedef void (*cb_t)(int);    /* a typedef tames the syntax — use one in interviews */
cb_t on_rx = uart_rx_cb; on_rx(42);
```

Macros vs inline vs const: a #define is textual paste — no type check, and MAX(a++, b) increments twice (the canonical trap). static inline functions give the same zero call cost with types and a debugger that can step them; const objects have addresses and obey scope. Preprocessor party tricks worth recognizing: #x stringizes an argument, a##b pastes tokens, and every header needs an include guard.

memcpy vs memmove — and implement memcpy (both are reported verbatim):

```
void *my_memcpy(void *dst, const void *src, size_t n) {
    uint8_t *d = dst;
    const uint8_t *s = src;
    while (n--) { *d++ = *s++; }
    return dst;
}
/* memmove adds ONE idea: if dst overlaps src and sits above it, */
/* copy backwards so you never overwrite bytes you still need. */
/* memcpy is allowed to assume no overlap — that is why it is faster. */
```

Say unprompted: production memcpy copies word-at-a-time once both pointers reach alignment — tying this back to the padding story above is the senior move.

Sixty-second drills — write these cold:

```
a ^= b; b ^= a; a ^= b; /* swap without temp — trap: zeroes both if &a == &b */
for (c = 0; v; c++) v &= v - 1; /* count set bits (Kernighan): one loop per 1-bit */
uint16_t probe = 1;
is_little = *(uint8_t *)&probe; /* runtime check: LE puts LSB at low addr */
```

PART III • INSIDE THE MICROCONTROLLER

What the silicon is, which chip to pick, how it boots — and the first peripherals every firmware touches.

9 MCU Width: 8-bit vs 32-bit

"Width" is the size of the CPU registers and ALU (arithmetic-logic unit — the part that actually adds and compares) — the chunk of data processed per instruction. An 8-bit AVR adds two 32-bit numbers in ~4+ instructions; a Cortex-M (ARM's microcontroller CPU family — this course's default brain, \$10) does it in one.

Why? Why does 32-bit dominate now? One instruction moves 4 bytes, addresses gigabytes linearly, and runs an RTOS + comms stacks comfortably. Modern 32-bit parts (Cortex-M0) even undercut 8-bit on price.

→ **Why that?** Why do 8-bit parts still ship in billions? Tiny die, tiny package, μ A sleep currents, 5 V tolerance, decades of proven silicon — a thermostat or toy needs none of the 32-bit muscle.

→ **Why that?** Why not "always use 32-bit to be safe"? Because engineering is fit-for-purpose: wider buses and bigger flash cost power, board area and money. Over-specifying is a design smell interviewers probe for.

Q: Why not always use a 32-bit MCU?

A. It can waste power and cost on simple tasks ✓ *The core reason — pay for what the product needs, not for spec-sheet comfort.*

B. It can be needless complexity for trivial tasks ✓ *Also valid: bigger toolchains, clock trees, more configuration surface for a job an 8-pin part does.*

C. 8-bit is always cheaper ✗ *Feels intuitive ("smaller = cheaper") but volume Cortex-M0 parts now beat many 8-bit prices — never say "always" in an interview.*

D. 32-bit is always faster ✗ *Usually, yes — but "always" dies on interrupt latency corner cases and clock-speed comparisons. Absolutes are the trap here.*

Interview tip: Two answers are right on purpose — real interviews use "select all that apply" to catch people who stop at the first plausible option.

10 CISC vs RISC

	RISC (ARM, RISC-V)	CISC (x86)
Instructions	Few, simple, fixed-length	Many, complex, variable-length
Work per instruction	Small (load/store separated from ALU ops)	Large (one instr. may read memory, compute, write back)
Pipelining	Easy — every stage takes similar time	Hard — instructions decode to internal μ -ops first
Silicon spent on	Registers, low power	Decoders, compatibility
Home turf	Embedded, mobile, now servers too	Desktop/server legacy software

Why? Why does RISC own embedded? Fixed-length simple instructions → small decoder → less silicon → less power and cost — the three currencies of embedded.


→ **Why that?** Why does easy pipelining follow from simple instructions? A pipeline wants every stage equal-length; one 20-cycle CISC monster instruction stalls everything behind it. Uniform instructions keep the conveyor belt full.

→ **Why that?** Why hasn't CISC died? Decades of x86 binaries. Ironically, modern x86 chips translate CISC to RISC-like μ -ops internally — RISC won the argument even inside its rival.

Q: Why is RISC preferred for embedded systems?

- A. Lower power consumption ✓ *Smaller decode logic and simpler datapaths directly cut dynamic and leakage power.*
- B. Easier pipelining ✓ *Uniform, fixed-length instructions keep pipeline stages balanced.*
- C. RISC is always faster ✗ *Sounds like the point of RISC — but per-task speed depends on compiler, memory, clock. A CISC chip can outrun a RISC one; "always" is the giveaway.*
- D. RISC is always cheaper ✗ *Usually true at the low end, but price is set by volume, process node, and licensing — not the ISA alone.*

11 Embedded Architectures



The layered view — every architecture question is really "which layer are we talking about?"

System types: standalone (washing machine), networked (smart meter), mobile (drone), real-time (ABS brakes — deadlines are part of correctness).

Firmware styles: **super-loop** (`while(1){read; compute; write;}` — simple, fully deterministic by inspection), **layered/event-driven** (interrupt-fed state machines), **RTOS** (preemptive tasks with priorities).


CPU ISAs (instruction-set architectures — the vocabulary a CPU understands): AVR/PIC (8-bit, Harvard, hobby-to-appliance), **ARM Cortex-M** (32-bit RISC, the industry default — M0 for cost, M4/M7 add DSP+FPU for control loops), **RISC-V** (open ISA, no license fee, exploding in new silicon).

Why? Why does architecture choice matter more than raw specs? It fixes your ecosystem: toolchain, RTOS ports, debugger support, hiring pool, and 15-year part availability. You can swap a resistor; you cannot cheaply swap an ISA mid-product.

→ **Why that?** Why is super-loop still respectable? For safety review, "one loop, no preemption" is trivially analyzable — some certified designs prefer it precisely because nothing can interleave.

12 Power-On to main() — the Boot Story

"What happens when the board powers on?" is the classic depth probe — it walks through the vector table, the linker script, and the \$7 memory map in one question.



The whole boot on one path — a bootloader is just a detour ending in the same jump.

Why? Why does a Cortex-M start by reading **data**, not executing code? Address 0 holds the vector table: word 0 is the **initial stack pointer**, word 1 the **reset handler address**. Hardware loads SP first because the very first C function or fault handler needs a working stack — the chip bootstraps the one resource everything else assumes.

→ **Why that?** Why can't reset jump straight to `main()`? Because `main()` assumes the C world already exists: initialized globals hold their values (`.data`, copied flash→RAM by startup code) and uninitialized ones are zero (`.bss`, zeroed in a loop). Startup code builds that world — copy, zero, set clocks/FPU — then calls `main`. Skip the copy and every initialized global is garbage: a bug candidates are asked to diagnose.

→ **Why that?** Why does the **linker script** keep coming up? It is the map-maker: it decides that `.text` lives in flash, where `.data`'s flash image and RAM home are, where the stack starts — the addresses the startup code and §7's memory map both obey. "Who places variables at addresses?" — the linker, directed by that script.

Bootloaders & safe OTA — the follow-up in any connected-product interview: a bootloader is a small resident program that runs first, decides "boot the app or update it", and jumps by loading the app's own vector table (set VTOR — the register that tells the CPU where the vector table now lives — then load the app's SP and reset vector). Field-safe updates are a checklist: **two slots** (never erase the only working image), **verify before boot** (a CRC — a math fingerprint that catches corruption — or a cryptographic signature that also proves who built it), and a **watchdog + rollback flag** so a bad image that boots and hangs gets demoted automatically. Anti-bricking is the whole design, not a feature.

One-liner to have ready: JTAG vs SWD — same ARM debug access, SWD does it over 2 pins instead of 4–5.

What the vector table actually looks like in a startup file — §8's array of function pointers, placed at address 0 by the linker:

```
__attribute__((section(".isr_vector")))
void (* const vector_table[])(void) = {
    (void (*)(void))(&_estack),      /* word 0: initial stack pointer      */
    Reset_Handler,                  /* word 1: where execution begins     */
    NMI_Handler, HardFault_Handler, /* word 2, 3, ... one per exception  */
};
```

Q: On a Cortex-M, what does the very first word of the vector table (address 0) hold?

- A. The first instruction to execute **✗** *True for many older CPUs (x86, AVR jump vectors) — which is exactly why it is offered. Cortex-M fetches data first.*
- B. The address of the reset handler **✗** *One slot off: that is word 1. Tempting because it is the second thing hardware reads.*
- C. The initial stack pointer value **✓** *SP before PC — so even the reset handler could be written in C and any early fault has a stack to run on.*
- D. The interrupt enable register **✗** *Confuses the vector table (addresses in flash) with NVIC peripheral registers (control bits in the peripheral space).*

"What happens from power-on to `main()`?" — the 45-second spoken answer

Reset releases and the core reads the vector table: word zero becomes the initial stack pointer, word one is the reset handler, and execution jumps there. The startup code then builds the world C assumes — it copies the `.data` section's initial values from flash into RAM, zeroes `.bss`, brings up clocks and the FPU, runs any C runtime init, and finally calls `main()`. The linker script chose all those addresses. And that vector-table mechanism is exactly what a bootloader exploits: to launch an application it just points VTOR at the app's table, loads its stack pointer and reset vector, and jumps.

“How does an interrupt work, end to end?” — the 45-second spoken answer

A peripheral raises its line into the NVIC, which compares priority against what is running. If it wins, the core finishes the current instruction, and hardware automatically pushes R0-R3, R12, LR, PC and xPSR onto the stack — the caller-saved set — then fetches the handler address from the vector table, which is literally an array of function pointers, and enters the ISR. On exit, a magic EXC_RETURN value in LR tells the hardware to unstack and resume. Back-to-back interrupts tail-chain without re-stacking, and total entry latency is about 12 cycles — plus however long your longest critical section keeps interrupts masked, which is the part you control.

13 Platform Face-off: Arduino vs STM32 vs ESP32

First, defuse the category error: **Arduino is a board family plus a software framework, not a chip architecture** — the Arduino core runs on STM32 and ESP32 too. The real comparison is between the silicon those names usually stand for: the 8-bit AVR ATmega328P, ST’s Cortex-M line, and Espressif’s radio-integrated SoC.

	ATmega328P (Arduino Uno)	STM32 (F4/G4 class)	ESP32
Core	8-bit AVR, Harvard (separate roads for code and data)	32-bit ARM Cortex-M4F with FPU (hardware float math — no slow software emulation)	Dual-core 32-bit Xtensa LX6; C3/C6 variants use RISC-V — \$10’s open ISA shipping in volume
Clock	16 MHz	168–170 MHz	240 MHz × 2 cores
RAM / Flash	2 KB / 32 KB	128 KB–1 MB / up to 1 MB+	520 KB / 4–16 MB external
Signal peripherals	Basic timers, 10-bit ADC, one each UART/SPI/I2C	Advanced motor timers (dead-time, complementary outputs), 12-bit ADCs with injected triggers, CAN, USB, DMA everywhere	Solid set incl. MCPWM — but the ADC is famously non-linear, and strapping pins constrain the GPIO map
Radio	None	None on-chip (add a PHY or module)	Wi-Fi + Bluetooth built in , module pre-certified
Toolchain	Arduino IDE / avr-gcc	CubeIDE + HAL/LL; first-class Zephyr and FreeRTOS ports	ESP-IDF (FreeRTOS-based) or the Arduino core
RTOS story	None — super-loop territory	Optional, your choice	FreeRTOS effectively mandatory : the radio stack needs it
Debug	UART bootloader, no on-chip debug	SWD: single-step, breakpoints, trace	JTAG (shares pins); USB debug on newer parts
Power	µA sleep achievable with care	µA stop modes, fine-grained clock gating	µA deep sleep — but 100 mA-class bursts whenever the radio speaks
Sweet spot	Teaching, one-evening prototypes, cost-floor gadgets	Industrial products, motor control, anything certified or long-lived	Connected devices where integrated Wi-Fi/BLE outweighs analog quality

Why? Why does ESP-IDF ship an RTOS whether you like it or not? The Wi-Fi and Bluetooth stacks are large concurrent state machines with real deadlines. Espressif pins them to one core (PRO_CPU) and hands you the other (APP_CPU) — the dual core exists largely so the radio never fights your application.

→ **Why that?** Why does STM32 own industry? Portfolio breadth (one toolchain from M0 to M7), 10-year longevity commitments, safety-certified lines, and analog you can trust — companies buy the *roadmap*, not the chip. When asked “why STM32 over ESP32 for this product?”, roadmap + ADC quality + power ceiling is the expected shape of the answer.

→ **Why that?** Why is “just use ESP32, the Wi-Fi is free” a trap? The pre-certified module is a huge win — but you inherit a 100 mA-class radio power budget, a non-linear ADC, strapping-pin constraints, and a single-vendor supply chain. Free features are never free; saying that sentence out loud is the senior signal.

Q: You are building a 20 kHz field-oriented (FOC) motor controller. Which platform?

- A. Arduino Uno** — simplest to program ❌ *The prototyping reflex. 16 MHz, no FPU for the Park/Clarke math, and no complementary PWM pairs with hardware dead-time — the one feature FOC cannot live without.*
- B. ESP32** — 240 MHz dual-core beats them all ❌ *The MHz trap. MCPWM exists, but the non-linear ADC poisons current sensing and radio/RTOS activity injects jitter into a loop that tolerates none. Possible / the fit.*
- C. STM32 G4/F4** — advanced timers + triggered ADC + FPU ✅ *TIM1/TIM8 give dead-time and center-aligned PWM in silicon; injected ADC conversions fire synchronized to the PWM valley; the FPU (plus CORDIC on G4) eats the trig. This family was shaped by this exact job.*
- D. Any of them** — it is just software ❌ *The answer that ends hardware interviews. Dead-time insertion and PWM-synchronized sampling are silicon features; software cannot conjure them.*


14 Timers: Hardware vs Software

	Hardware timer	Software timer
What it is	Dedicated counter peripheral clocked from the MCU clock via a prescaler (a divide-by-N stage that slows the raw clock to a useful rate)	A counter variable decremented on each tick of SysTick (the core's built-in heartbeat timer) — FreeRTOS software timers, Zephyr k_timer
Precision	Cycle-accurate (ns-μs), zero jitter	Tick-limited (typically 1 ms) plus scheduler jitter
CPU cost	None while counting; CPU only services the interrupt	CPU walks the timer list on every tick
Quantity	Fixed by silicon (e.g. 4-14)	Practically unlimited
Use it for	PWM, motor commutation, input capture, encoder counting	Timeouts, blinking LEDs, periodic housekeeping tasks

Why? Why does hardware win on precision? The counter increments in silicon every clock edge — no instruction fetch, no scheduler in the path. A software timer cannot see time finer than its tick.

→ **Why that?** Why not make the RTOS tick 1 μs and get precise software timers? Then the tick ISR fires a million times per second and eats the CPU. Tick rate is a trade: resolution vs overhead — which is exactly why both kinds exist.

Worked example — 20 kHz PWM from a 72 MHz clock: pick $(PSC+1) \times (ARR+1) = 72 \text{ MHz} / 20 \text{ kHz} = 3600 \rightarrow PSC = 0, ARR = 3599$; duty = $CCR/(ARR+1)$, so $CCR = 900$ gives 25%. The resolution falls out for free: 3600 steps ≈ 11.8 bits — and raising the PWM frequency trades that resolution away, a follow-up interviewers enjoy.



One clock source, two paths: silicon precision for signals, tick-based flexibility for scheduling.

Q: You need a 20 kHz motor PWM with <100 ns edge accuracy. Which timer?

- A. RTOS software timer ❌ *Feels fine because "20 kHz is fast, RTOS is fast" — but a 1 ms tick cannot even represent a 50 μ s period, let alone 100 ns edges.*
- B. Hardware timer in PWM mode ✅ *Edges are generated in silicon from the prescaled clock; the CPU is not even involved per edge.*
- C. A delay loop (busy-wait) ❌ *Tempting for its simplicity — accurate-ish, but burns 100% CPU and any interrupt shatters the timing.*
- D. SysTick ❌ *Sounds "hardware" — it is, but it exists to drive the OS tick, is 24-bit, and has one shared interrupt, not PWM outputs.*


15 Watchdog Timer

A **watchdog (WDT)** is a hardware down-counter that resets the MCU unless firmware "feeds" (reloads) it before it hits zero. Feeding is proof of life; silence is treated as a crash.

Why? Software in the field *will* eventually hang — cosmic-ray bit flips, deadlocks, corrupted pointers, a stuck I2C bus. The watchdog turns "frozen forever" into "down for 100 ms".

→ **Why that?** Why must it be **hardware**, not a software check? Because a hung CPU cannot run the software that checks whether the CPU is hung. The guard must live outside the thing it guards — ideally on an independent clock, so a crystal failure cannot kill both.

→ **Why that?** Why feed from the main loop, not from a timer ISR? Timer interrupts often keep firing even when main code is deadlocked — feeding from an ISR makes the watchdog blind. Golden rule: **feed only where real work provably completed.**



Watchdog: healthy feed cycle (top) vs missed feed forcing a hardware reset (bottom).

Q: Why use a watchdog timer?

- A. To increase CPU speed ❌ *Tempting because timers relate to clocks — but the WDT never accelerates anything; it only counts down to a reset.*
- B. To reset the MCU if software hangs ✅ *Its one job: autonomous recovery from a hung system.*
- C. To save power ❌ *Sounds plausible since sleep modes use timers to wake up — but that is the RTC/low-power timer, a different peripheral.*
- D. To reduce memory usage ❌ *Pure distractor; watchdogs never touch memory footprint.*


Interview tip: Follow-up they love: "What is a **windowed** watchdog?" — feeding too early also resets, catching runaway loops that spam the feed instruction.

Numbers to have ready — STM32 IWDG: clocked from the ~32 kHz LSI, prescaler 32 → 1 kHz count rate, reload 2000 → 2 s timeout, so feed at ≤ 1 s from the main loop. And say the catch: the LSI is loosely specified across temperature, so a watchdog budget needs margin, never precision.

16 The Peripheral Zoo — ADC, DAC, PWM & Friends

A **peripheral** is a built-in helper gadget sitting beside the CPU — a specialist circuit that does one hardware job on its own. Here are the employees every interview assumes you have met — what each does, in plain words, with

the numbers and the trap that comes attached. They all serve one loop:



the whole embedded job in one line: measure → decide → act. (Swap PWM for a DAC when you need a smooth voltage instead of switching.)

Every embedded product is some version of this chain — the peripherals are its specialist workers.

GPIO — the simplest employee. One pin, read or driven (§3’s register lines). Its *modes* are the interview material: input with an internal pull-up/down (so an unconnected pin is never floating and random); output push-pull (actively drives high AND low — the default); output open-drain (can only pull low — the sharing trick §22’s I2C is built on); and alternate function, which hands the pin over to another peripheral entirely — the same metal leg becomes the UART’s TX or the timer’s PWM output.

ADC — the measuring tape. It converts a voltage into a number by laying 2ⁿ evenly spaced marks between 0 and V_{ref} and reporting which mark the input is nearest. So each mark (one LSB) is V_{ref} / 2ⁿ: a 12-bit ADC at 3.3 V resolves steps of 0.81 mV, and a raw reading of 2048 means 2048/4096 × 3.3 = 1.650 V. Two rules travel with it: sample faster than 2× the fastest frequency *present* and filter before converting (§39 Nyquist), and never confuse resolution (step size) with accuracy (how true the steps are — §13’s ESP32 complaint).

Why? Why does every ADC start with a **sample-and-hold** capacitor? You cannot measure a moving target — the capacitor takes a snapshot of the input and holds it perfectly still for the whole conversion. Photograph first, measure the photograph.

→ **Why that?** Why does a high-impedance sensor break it? That capacitor charges *through your source resistance*. Feed it from 1 MΩ and it cannot reach the input voltage within the sample window — it still partly remembers whatever it measured last. Fix: a buffer op-amp in front, or configure a longer sampling time.

→ **Why that?** Why does averaging 16 readings feel like extra bits? Noise is random, signal is not — averaging N samples shrinks the noise by √N, buying roughly one extra clean bit per 4× oversampling. Free resolution, paid for in time.

Q: A 12-bit ADC scans several channels; the channel fed by a 1 MΩ sensor reads like a blend of itself and the previously scanned channel. Cause?

- A. The ADC is faulty **✗** *The blame-the-part reflex — but a fault would not produce so specific a symptom as “mixed with the previous channel”.*
- B. V_{ref} is unstable **✗** *A wobbling reference rescales every channel’s reading; it cannot blend one channel into another.*
- C. The sample-and-hold capacitor cannot charge through 1 MΩ in the sample window — it still remembers the last channel **✓** *The classic multiplexed-ADC ghost. Fix: buffer the sensor with an op-amp, or lengthen the sampling time.*
- D. The result variable is missing volatile **✗** *§6 imported into an analog problem — software visibility cannot mix voltages.*

“How does an ADC work?” (SAR — the common kind) — the 45-second spoken answer

First, a sample-and-hold capacitor freezes the input — you cannot measure a moving target. Then it plays the number-guessing game: an internal DAC proposes half of V_{ref} , a comparator answers higher-or-lower, and that answer becomes the top bit; halve the remaining range and repeat for every bit. N comparisons give N bits — which is exactly why conversion takes about one clock per bit. The result is a count of steps of size V_{ref} over 2^n . And the honest caveat: that resolution number hides the three things accuracy actually depends on — the reference, the linearity, and giving the capacitor time to charge through your source impedance.

DAC — the tape run backwards. A number in, a real voltage out (classically via an R-2R resistor ladder — each bit switches a rung, the rungs sum). Used for audio, waveform generation, and analog control signals. Its trap: the output is *weak* — it sets a voltage but cannot supply real current, so anything hungrier than an op-amp input gets a buffer. Many small MCUs simply omit the DAC, because the next peripheral fakes one well enough...

PWM — digital pretending to be analog. The pin is only ever fully ON or fully OFF, but it flips so fast that the world downstream sees only the average: $V_{avg} = \text{duty} \times V_{cc}$. A dimmed LED, a motor at half speed, and (through an RC filter) a stand-in DAC are all the same trick at different speeds. It is made by a timer — §14’s worked example (72 MHz $\rightarrow$ 20 kHz, 3600 steps) is literally this — and it carries §14’s trade: more frequency, fewer distinct duty steps. One special dialect worth naming: hobby-servo PWM, 50 Hz where only the 1–2 ms pulse *width* encodes the commanded angle.

Why? Why PWM a motor instead of feeding it a lower voltage? A transistor dropping voltage *linearly* burns the difference as heat: $P = (V_{cc} - V_{motor}) \times I$, watts of it. A switching transistor is either fully on (near-zero volts across it) or fully off (zero current through it) — either way, almost zero waste. §18’s whole gate-drive obsession exists to keep the switch out of the in-between region where it briefly behaves like the wasteful linear version.

Timers beyond PWM. §14 covered the counting; two timer *modes* earn their own names in interviews: input capture latches the timer’s count in hardware the instant an edge arrives — frequency and pulse-width measurement with zero software jitter; and encoder mode lets the timer do §34’s $\times 4$ quadrature counting entirely in silicon.

DMA — the courier. A peripheral that exists to move data between other peripherals and RAM *without the CPU carrying each byte*: point it at the ADC, hand it a buffer, and samples stream into memory while the CPU does the control math — the dashed path in the diagram. Hire a mover instead of carrying one box per trip. Its interview trap lives in §39: on cached cores, the CPU may not *see* what DMA just wrote until coherency is handled.

Honourable mentions, one line each:


Peripheral	What it is, in plain words
Comparator	An instant analog referee: output flips the moment one input crosses the other — threshold detection with no sampling, no software, and it can fire an interrupt.
RTC	A wall-clock that keeps time through sleep and even power-off (coin cell) — seconds and calendars, not microseconds.
CRC engine	Hardware that computes a CRC (the mathematical fingerprint frames carry to detect damage) in one instruction per word instead of a software loop.
Temperature sensor	A built-in thermometer on the die — coarse, but free; good for “is the chip cooking?”, not lab metrology.
Watchdog	Already has its own section — §15: the dead-man’s switch.

Peripheral	What it is, in plain words
UART / SPI / I2C / CAN / USB blocks	The protocol engines: each one is a peripheral that speaks one of Part V's languages in hardware, so the CPU only loads and unloads mailboxes.

PART IV • THE CIRCUITS FIRMWARE TOUCHES

Enough hardware judgment to be trusted near a schematic.

17 Resistor Networks: the Delta Question



Seen from A and B, there are only two paths:

direct: A-B via R3 = 30

the detour: A-C-B via R1+R2 = 60

$$R_{AB} = 30 \parallel 60 = 20$$

no transform needed — just see the two paths

Equal delta, R_{AB} asked. The interviewer is testing whether you reach for a formula or for your eyes.

Why? Why is this a trick question at all? Because candidates who memorized the Delta-Y transform apply it, get Y legs of $R/3 = 10$ each, and then **forget that A-to-B passes through two Y legs** ($10 + 10 = 20$). The transform works — the half-remembered transform is the trap.

→ **Why that?** Why does the two-path view always win here? From A to B the network is literally R3 in parallel with $(R1 + R2)$: $30 \parallel 60 = (30 \cdot 60) / 90 = 20$. Ten seconds, no formula, no place to slip. Reach for transforms only when no two nodes see a clean series/parallel structure (e.g. bridge circuits).

Q: Delta of three 30 Ω resistors — what is R_{AB} ?

A. 90 Ω — all three in series **✗** The "count the resistors" reflex. A and B are directly joined by R3; nothing is in plain series.

B. 10 Ω — Delta-to-Y gives 10 Ω per leg, so $R_{AB} = 10 \Omega$ **✗** The dangerous one: the transform step is correct, the conclusion isn't — A to B crosses two Y legs: $10 + 10 = 20 \Omega$.

C. 20 Ω — R3 (30) in parallel with R1+R2 (60) **✓** The two-path view. $30 \parallel 60 = 20 \Omega$.


D. 15 Ω — all three in parallel, doubled **✗** Invented arithmetic dressed in confident language — a reminder that plausible-sounding / derivable.

18 MOSFET Switching: High-Side, Low-Side & Gate Drive


The vocabulary first: a **MOSFET** is the standard electronic switch. Three legs — the **gate** is the control knob; **drain** and **source** are the two ends of the switch itself. **V_{gs}** means the voltage measured from gate to source, and the switch turns on once V_{gs} exceeds a threshold **V_{th}** (a few volts). NMOS and PMOS are its two flavors: N switches on with a positive V_{gs}, P with a negative one. With that vocabulary, the trap:

The single most-failed screening question in hardware rounds. Everything hinges on one fact: **V_{gs} — gate voltage measured relative to the source — is what turns a MOSFET on**, not gate voltage relative to ground.

Low-side switch — works



High-side NMOS, 3.3 V gate — does not turn on



Same transistor, same 3.3 V GPIO — topology decides whether V_{gs} ever exceeds V_{th} .

Why? Why does the high-side NMOS fail? The source follows the load node. The instant the FET starts conducting, the source rises toward 24 V — which erases V_{gs} . The FET throttles itself into a half-on equilibrium (source floats a few volts below the gate) and cooks in the linear region, or never turns on at all.

→ **Why that?** Why not just fix it with a resistor or a bigger GPIO? Because it is not a drive-strength problem, it is a **reference** problem. Fixes must move the gate *above the supply*: a bootstrap gate driver (flying capacitor — needs continuous switching to stay charged), a charge pump (works for DC-on), or side-step it entirely with a **P-channel** FET, whose gate is pulled *below* its source (the 24 V rail) to turn on.

→ **Why that?** Why do NMOS-everywhere designs exist anyway? Electron mobility → NMOS has ~half the $R_{ds(on)}$ (residual resistance when fully on) of an equal-size PMOS — so power stages (motor bridges) pay for bootstrap drivers to keep NMOS on the high side. Low-current loads just use a PMOS.

Q: N-channel MOSFET, source referenced to a 24 V rail, load on the other side, gate driven by a 3.3 V GPIO. Does it switch?

A. Yes — $V_{gs} = 3.3$ V is above V_{th} , so it turns on **X** *The trap for anyone measuring gate voltage from ground. V_{gs} is gate-to-source, and the source is at 24 V — V_{gs} is hugely negative.*

B. No — the gate would need ~ 24 V + V_{th} ; use a bootstrap/charge-pump driver or a P-channel instead **✓** *The reference problem stated correctly, plus the three standard fixes.*

C. No — but only because the GPIO cannot source enough gate current; a driver IC is always required **X** *Half-true bait: GPIOs are weak drivers, but "always required" is false (low-side, slow switching is fine) — and current is not this circuit's failure mode.*

D. Yes — provided a 10 k gate pull-down is added so it turns off cleanly **X** *Pull-downs are genuinely good hygiene (defined state during MCU reset) — which is what makes this tempting. But hygiene cannot conjure 27 V.*

Gate drive quality — the fall-time question. A gate is a capacitor (~nF). Edge speed is set by how fast that capacitance charges/discharges: $t \approx R_{gate} \cdot C_{gate}$. Slow edges leave the FET in the linear region where it dissipates $V \cdot I$ as heat every single switching cycle.

Why? Why is a slow *fall* specifically common? Turn-off discharges the gate through the GPIO/driver's pull-down path plus the series gate resistor — GPIO pull-downs are weak, and the Miller plateau (a capacitive-feedback stall: the swinging drain pumps charge back into the gate, freezing its voltage mid-transition) holds the gate hostage.

The argument in numbers: a 20 nC gate discharged by a 0.5 A driver takes $t \approx Qg/I = 40$ ns; through a GPIO's ~20 mA it takes 1 μ s — 25× longer in the linear region, every cycle, 20 000 times a second. That one division is the entire "why a gate driver" case.

Q: Gate fall time is much slower than rise time, causing turn-off switching losses. Which actions directly reduce fall time? (select all that apply)

- A. Decrease the gate resistor value** ✓ *Smaller $R \rightarrow$ faster discharge of the same capacitance. Direct and correct (watch ringing).*
- B. Use a dedicated gate driver IC** ✓ *Drivers sink amps into the gate — orders of magnitude stronger pull-down than a GPIO.*
- C. Put a diode anti-parallel with the gate resistor so turn-off current bypasses it** ✓ *Classic asymmetric-drive trick: R limits the rise, the diode shortcuts the fall.*
- D. Increase the PWM frequency** ✗ *Feels related because "switching" — but edge speed is an RC property, frequency-independent. More cycles per second just multiplies the loss.*
- E. Raise the gate-driver supply voltage** ✗ *Speeds the rise (more overdrive), but the fall still discharges through the same R to 0 V — starting from higher voltage, so if anything the fall to V_{th} takes longer.*

PART V · HOW MACHINES TALK

One protocol at a time — from two bare wires to publish/subscribe over the network.

19 The Protocol Map — Pick by Constraint

Protocol	Wires	Speed	Killer feature	Watch out
UART	2 (TX/RX)	~115 kbps typ.	Dead simple; every chip has one; the debug console	No clock — both ends must agree on baud; point-to-point only
SPI	4 (+CS each)	10–100 Mbps	Fastest simple bus; full duplex	Pin count explodes with devices; no ACK — errors are silent
I2C	2 (SDA/SCL)	100k–3.4 Mbps	Many devices, 2 wires, addressing + ACK built in	Open-drain $\rightarrow$ needs pull-ups; a stuck slave can hang the bus
CAN	2 (diff pair)	1 Mbps (8 Mbps FD)	Multi-master, priority arbitration, brutal noise immunity — automotive DNA	Payload only 8 bytes (64 in FD); needs a transceiver (the small chip converting the MCU's logic-level TX/RX into the bus's differential voltages) + termination
Ethernet	4–8	100 M–1 Gbps+	Bandwidth, IP ecosystem, ROS-native	Power, connector cost, non-deterministic unless TSN/EtherCAT
BLE / Wi-Fi / LoRa / Zigbee	radio	kbps–Mbps	BLE: phone+coin cell · Wi-Fi: bandwidth · LoRa: km-range at μW · Zigbee: mesh	Pick by the range–power–bandwidth triangle: you get two of three
MQTT / HTTP / CoAP	app layer	—	MQTT: pub/sub for flaky links · CoAP: REST for tiny nodes	These ride on TCP/UDP — they are not alternatives to the buses above

Interview move: place any protocol on this stack before comparing it.

Application	MQTT · HTTP · CoAP · DDS (ROS 2)
Transport / Network	TCP/UDP · IP · CANopen · 6LoWPAN
Data link	CAN frames · Ethernet MAC · BLE link layer
Physical	UART · SPI · I2C wires · CAN diff pair · radio


Q: Ten sensors on a robot arm, 2 spare MCU pins, moderate speed. Which bus?

- A. SPI **✗** The "fast = good" reflex — but 10 chip-selects need ~13 pins. You have 2.
B. I2C **✓** Two wires, 10 addresses, ACKs for error detection — exactly its design point.
C. UART **✗** Feels universal, but it is point-to-point: 10 sensors would need 10 UARTs.
D. CAN **✗** Would work electrically and is robust — tempting! But it needs a transceiver per node and protocol overhead that plain sensor reads don't justify. Over-engineering is also a wrong answer.

Each of the next seven sections opens one row of this map to interview depth — read them in order: the wires get smarter as you go.

20 UART — Two Wires and an Agreement

Asynchronous, point-to-point, no shared clock: TX of one device to RX of the other, and both sides **agree on the baud rate in advance**. The frame does the rest:



shown: 0x41 sent LSB-first over 8N1 — 1 start + 8 data + 1 stop = 10 bits per byte (so 115200 baud ≈ 11.5 kB/s)

One 8N1 frame. The receiver samples each bit at its center, typically with 16× oversampling.

Why? Why does clockless work at all? Every byte starts with a falling START edge — the receiver restarts its bit timing there, so clock error only has to stay small across **10 bits**, not forever. That bounds the tolerance: total drift must stay under about half a bit in 10, roughly ±2–3% baud mismatch. It also explains a classic field bug: an MCU on a sloppy internal RC oscillator drifts with temperature until 115200 starts framing-erroring.

→ **Why that?** Why is there a STOP bit if it is just idle-high? It guarantees a high level before the next START's falling edge — no stop, no edge, no resync. It is the minimum gap that keeps the scheme self-timing.

Numbers worth memorizing: at 115200 baud, one byte = 10 bits ≈ 87 μs — so printf-ing a 50-character line costs ~4.3 ms. A blocking print inside an ISR at that price is how “adding a debug log changed the bug” stories are born.

Flavor	Levels	Reach / topology	Where you meet it
TTL / LVTTTL UART	0 V and 3.3 V (or 5 V), idle high	Centimeters; board-level, point-to-point	MCU ↔ MCU, debug consoles, GPS modules
RS-232	±3...±15 V, inverted logic	~15 m, point-to-point	Legacy industrial gear, PC serial — never wire it straight to a 3.3 V pin
RS-485	Differential pair (CAN's trick, §23)	Up to ~1200 m, multi-drop half-duplex bus	Modbus RTU, industrial sensor strings; needs direction (DE/RE) control

Add hardware flow control (RTS/CTS) when the receiver can be slower than the stream — it is a “stop sending, buffer full” wire, nothing more. And 8N1 decodes as: 8 data bits, No parity, 1 stop bit.

Q: Board A transmits at 115200 baud; board B is misconfigured to 9600. What does B receive?

A. Nothing at all — the UART ignores wrong-speed traffic **✗** *Half right, which is the bait: the hardware cannot know the speed is wrong. It happily samples — just in the wrong places.*

B. The data, but slower **✗** *Imports the TCP mindset. UART has no negotiation, no retry, no rate adaptation — nothing “adjusts”.*

C. Garbage bytes plus framing errors **✓** *B samples where it expects bits and finds mid-transition levels; STOP-bit checks fail → framing errors, and what passes decodes as junk.*

D. A damaged RX pin from the over-speed signal **✗** *Electrical-sounding, but speed is not voltage — the pin sees the same levels either way.*


Interview tip: Follow-up they love: “how does auto-baud work?” — measure the width of a known pattern (e.g. 0x55, alternating bits) and derive the divisor.

“How does UART work?” — the 45-second spoken answer

UART is asynchronous serial: TX of one device to RX of the other, no clock wire, both sides agreeing on a baud rate in advance. The line idles high; a falling START bit announces a byte and re-synchronizes the receiver’s timing. It then samples 8 data bits LSB-first at each bit center — typically with 16× oversampling — optionally checks parity, and expects a high STOP bit; a low there is a framing error. Because timing resets every byte, the two clocks only need to agree within a couple of percent. It is point-to-point and unacknowledged — simple and universal, which is why it is every chip’s debug port.

21 SPI — the Shift-Register Ring

Four wires — SCK, MOSI, MISO, CS — and one idea: master and slave are two halves of a single circular shift register (a row of bit-boxes where every clock tick pushes each bit one slot along).



one 16-bit ring: every SCK edge rotates it one step — a byte OUT is always a byte IN

Full duplex is not a feature, it is the mechanism: bits rotate around the ring, one per clock.

Why? Why must you transmit dummy bytes to *read*? Only the master generates SCK, and MISO only shifts on clock edges. No transmission → no clocks → the slave’s answer stays frozen in its register. “Send 0xFF to receive” is not a hack — it is the ring doing its only trick.

→ **Why that?** Why four “modes”? Two free choices — clock idle level (CPOL) and which edge samples (CPHA) — were made differently by different vendors, so four combinations exist. Mismatch the mode and everything arrives bit-shifted or edge-inverted. Mode 0 is the most common; the datasheet phrase to hunt for is “data valid on rising edge”.

→ **Why that?** Why does SPI have no speed ceiling in a spec? It is **source-synchronous**: the clock travels with the data, so there is no pre-agreed baud to honor (UART’s burden). The limit is physics — trace length and slave timing — which is how 50+ MHz flash chips share the same four wires as a 1 MHz sensor.

Mode	CPOL (idle level)	CPHA	Data sampled on
0	0 (low)	0	Rising edge — the default guess
1	0 (low)	1	Falling edge
2	1 (high)	0	Falling edge

Mode	CPOL (idle level)	CPHA	Data sampled on
3	1 (high)	1	Rising edge

No addressing (CS picks the chip), no ACK, no error detection — SPI is fast because it is **blind**. The professional reflex: after init, read a known register (WHO_AM_I) to prove the link, because SPI itself will never tell you it is broken.

Example — the first transaction of every IMU driver (MPU-series over SPI, mode 0):

```
cs_low();
spi_xfer(0x75 | 0x80); /* WHO_AM_I register; bit 7 set = READ */
id = spi_xfer(0xFF); /* dummy byte – its only job is making clocks */
cs_high(); /* id == 0x70 → wiring, mode and chip all OK */
```

Q: Why do SPI sensor drivers transmit dummy bytes during a register read?


- A. To give the slave processing time ✗ Borrowed from I2C clock-stretching intuition — but SPI slaves have no way to ask for time, and none is granted.
- B. The bus requires an even byte count ✗ An invented rule with a plausible ring to it. No such constraint exists.
- C. MISO only shifts when the master clocks — transmitting is how clocks get generated ✓ The ring mechanism itself. The dummy byte's content is irrelevant; its clocks are everything.
- D. The dummy bytes power the slave ✗ Parasite power exists — on 1-Wire, not SPI. Right concept, wrong bus: a classic cross-protocol distractor.

“How does SPI work?” — the 45-second spoken answer

SPI is a synchronous, master-driven bus. The master selects one slave with chip-select and generates the clock; master and slave each hold a shift register, wired into one ring through MOSI and MISO. Every clock edge rotates the ring one bit, so a byte sent is always a byte received — full duplex is the mechanism, not a feature, and reading requires transmitting dummy bytes just to generate clocks. The mode (CPOL/CPHA) fixes the clock idle level and sampling edge. No addressing, no ACK, no speed ceiling in the spec — it is fast because it is blind, so you prove the link yourself, typically by reading a WHO_AM_I register.

22 I2C — Many Devices, Two Wires, Built-in Manners

Two open-drain lines (SDA data, SCL clock) shared by every device; each chip owns a 7-bit address, and every byte is acknowledged. One complete write transaction:



control symbols live where data may not: SDA changing while SCL is HIGH = START/STOP.
Data must be stable during SCL-high — so those two “violations” are unmistakable, by design.

START, address + R/W, per-byte ACKs, STOP — the whole protocol in one line.

Why? Why open-drain? Devices may only pull LOW; the resistor supplies the high. Two chips “driving” at once can therefore never short-circuit — which is what makes a shared, even multi-master bus electrically safe. It also gives arbitration for free: a master that writes 1 but reads back 0 has lost and backs off — the same wired-AND trick CAN industrializes in §23.

→ **Why that?** Why is clock stretching a feature, not a bug? A slow slave may **hold SCL low** until it is ready, and the master must wait — flow control built into the physical layer. Also a classic field bug when a bit-banged or cut-corner master ignores it.

→ **Why that?** Why does **repeated START** exist? The universal sensor read is write-register-pointer-then-read. A new START without a STOP keeps the bus *yours* between the halves — a STOP would be the gap where another master hijacks the pointer you just set. Atomicity, bought with one symbol.

The hung-bus recovery ritual (interview favorite): if the master resets mid-read, a slave can sit holding SDA low forever, deadlocking the bus. Fix: toggle SCL up to 9 times so the slave finishes shifting its byte, then issue a STOP. Knowing this ritual = has actually debugged I2C.

Speeds: 100 kHz standard, 400 kHz fast, 1 MHz fast-plus, 3.4 MHz high-speed. 10-bit addressing exists but is rare.

Concrete transaction — burst-reading a BME280: START · 0x76+W · write register pointer 0xF7 · **repeated START** · 0x76+R · clock in 8 data bytes (ACK each, NACK the last to say “done”) · STOP. Every sensor datasheet’s “burst read” diagram is this exact shape — recognize it once, read them all.

Sizing the pull-ups — the design trade-off

Constraint	Sets which limit	The math
Rise time: line must reach a valid high within spec (1000 ns @ 100 kHz, 300 ns @ 400 kHz)	Maximum R	$t_{\text{rise}} \approx 0.85 \cdot R \cdot C_{\text{bus}}$. Small bus (~50 pF) at 100 kHz tolerates ~20 kΩ; a long 400 pF bus at 400 kHz needs ≤ 1 kΩ
Sink current: a device pulling low must reach $V_{\text{OL}} \leq 0.4$ V with ≤ 3 mA	Minimum R	$R_{\text{min}} = (V_{\text{dd}} - 0.4) / 3 \text{ mA} \approx 1 \text{ k}\Omega$ at 3.3 V
Battery life: every low period burns $I = V_{\text{dd}} / R$ through the resistor	Pushes R up	4.7 kΩ → 0.7 mA per held-low line; 10 kΩ → 0.33 mA — huge at μA sleep budgets

Why? Why is “4.7 kΩ, it’s the default” a weak interview answer? Because it skips the reasoning. Defaults encode *someone else’s* trade-off. Show the two limits, then place your value: fast bus → low R; battery node at 100 kHz → high R.

Q: Battery-powered sensor node, I2C at 100 kHz, short bus. Pull-up value?

- A. 1 kΩ** — fast edges, robust at all speeds ❌ *The signal-integrity reflex. Edges you don’t need, paid for at 3.3 mA per low line — a battery assassin.*
- B. 4.7 kΩ** — the standard default ❌ *Tempting precisely because it usually works. But “default” is not a reason, and here it leaves 2× battery life on the table.*
- C. 10 kΩ or higher** — slower edges are fine at 100 kHz, and static current halves or better ✅ *Both limits respected, trade-off argued from the product’s priority: battery life.*
- D. 100 Ω** — immune to bus capacitance ❌ *Violates the sink-current spec outright: devices can’t pull the line to a valid low. The bus stops working entirely.*

“How does I2C work?” — the 60-second spoken answer

I2C is a two-wire shared bus: SDA and SCL are open-drain with pull-up resistors, so devices only ever pull low and can never short each other — that is what makes many devices, even multiple masters, electrically safe on the same pair. The master issues a START — SDA falling while SCL is high — then a 7-bit address plus a read/write bit; the addressed slave ACKs by pulling SDA low on the ninth clock. Data moves byte by byte, each byte acknowledged, and a slow receiver can pace the master by clock stretching — holding SCL low. A repeated START chains a register-pointer write and the read atomically, and STOP — SDA rising while SCL high — releases the bus. 100 kHz standard, 400 kHz fast mode.

23 CAN — Arbitration, Frames & Fault Containment


Why the differential pair survives a panel full of motor drives. The receiver never looks at CANH or CANL alone — only at their **difference**. Noise coupling into the cable hits both wires nearly identically (common-mode), so it subtracts to zero at the receiver.

Why? Why does **twisting** matter, specifically? Cancellation only works if the noise really is identical on both wires. Twisting keeps the two conductors co-located and alternating in orientation along the whole run, so

every noise source couples equally into each — and the tiny loop area picks up almost no magnetic flux to begin with.

→ **Why that?** Why do the wrong explanations sound right? "It doubles the amplitude" — differential signalling *does* give 2× swing, but that is a signal-to-noise bonus, not the immunity mechanism. "The wires carry identical signals" — they carry **complementary** ones; if they were identical the difference would be zero and there'd be no signal at all.

Arbitration, worked with real IDs. Main controller ID 0x010 and end-effector ID 0x005 transmit at the same instant:




Bit-by-bit wired-AND race: first node to send recessive-1 while the bus reads dominant-0 backs off.

Why? Why is 0x005 the winner and not "the master"? CAN has **no master** — priority lives entirely in the ID. 0x005 has more leading zeros (dominant bits), so at bit 4 it holds the bus at 0 while 0x010 tried a 1. Lower numeric ID = higher priority, by construction.

→ **Why that?** Why is "collision + random backoff" (the Ethernet answer) wrong here? Because nothing was destroyed: the loser stopped *mid-frame* without corrupting the winner's bits. That is what "non-destructive" buys — the top-priority frame's worst-case latency is computable, which Ethernet CSMA/CD can never promise. It's the whole reason vehicles standardized on CAN.

ID allocation in practice: E-stop 0x001, motor commands 0x080–, sensor telemetry 0x200–, logging 0x500– — priority is a design-time spreadsheet, and the safety message wins every collision it will ever have, by construction.

Anatomy of a frame — and what each field buys



plus bit stuffing everywhere before the CRC delimiter: after 5 identical bits, a stuffed opposite bit is inserted


Why? Why bit stuffing? CAN is NRZ (non-return-to-zero: the wire simply holds each bit's level, so identical bits make a flat, edgeless line) — a long run of identical bits has no edges, and receivers resynchronize their sampling on edges. After 5 identical bits the transmitter **must** insert the opposite one (receivers strip it). Six identical bits in a row is therefore illegal — which is exactly what an error frame is: a deliberate rule violation every node instantly recognizes.

→ **Why that?** Why does the ACK slot work with *any* receiver asserting it? The transmitter sends the slot recessive; every node that received a valid frame pulls it dominant. One dominant bit proves "somebody heard me correctly" — cheap and anonymous — and its absence tells the transmitter to flag a TX error and retry.

→ **Why that?** Why the error-state ladder? Every node keeps error counters (TEC/REC). Healthy = **error-active** (may shout error frames); past 127 = **error-passive** (may only whisper, and waits extra before transmitting); past 255 = **bus-off** (silence until rehabilitation). A node with a failing transceiver thus **quarantines itself** instead of jamming the safety bus — fault containment as a first-class protocol feature, and the reason "what if a CAN node goes insane?" has the answer "the protocol already

thought of that.”

Termination: one $120\ \Omega$ resistor at **each physical end** of the bus — at 1 Mbps over meters of cable, an unterminated end reflects edges back and corrupts sampling. **CAN FD** keeps arbitration at the classic rate, then accelerates the data phase (up to ~8 Mbps) and stretches the payload to 64 bytes.



The error-state ladder: a sick node lowers its own voice before it can hurt the bus.

Q: Bus powered down, multimeter across CANH–CANL reads ~60 Ω . Diagnosis?

- A. Fault — a healthy bus should read $120\ \Omega$ ✗ The memorized-number trap: $120\ \Omega$ is one resistor, not the bus. Reading 120 actually means a terminator is missing.
- B. Healthy — two $120\ \Omega$ terminators in parallel measure $60\ \Omega$ ✓ The field technician’s one-second bus check — the answer that proves you have held the multimeter.
- C. Short circuit between the lines ✗ A short reads near $0\ \Omega$, not 60.
- D. A damaged transceiver loading the bus ✗ Transceivers are high-impedance when unpowered; they barely register on this measurement.

“How does CAN work?” — the 60-second spoken answer

CAN is a multi-master differential bus built for hostile environments. Any node may transmit; when two start simultaneously, bitwise arbitration on the identifier resolves it — a dominant 0 overwrites a recessive 1 on the wired-AND bus, so the lower ID wins without a single corrupted bit, and the loser retries when the bus frees. A frame carries the ID, up to 8 data bytes (64 in FD), a CRC, and an ACK slot that any receiver pulls dominant to confirm reception. Bit stuffing preserves edges for synchronization; error counters walk a faulty node from error-active to error-passive to bus-off, so sick nodes quarantine themselves; and the twisted differential pair cancels common-mode noise. That stack of guarantees is why vehicles and robots trust it with brakes.

24 Ethernet & TCP/IP — When the Robot Joins the Network

The moment data leaves one board, it gets wrapped in layers — each solving one problem, each with its own address:

each layer wraps the one above — switches read MACs, routers read IPs, endpoints read ports

Ethernet frame: dst MAC | src MAC | type | CRC

IP packet: src IP | dst IP | TTL | proto

TCP / UDP: src port | dst port

payload

Why? Why two addresses — MAC and IP? MAC (48-bit, factory-assigned) is for **local delivery**: switches forward frames by MAC within one segment. IP (routable, hierarchical) is for **global reach**: routers forward hop by hop across networks. ARP is the glue that answers “which MAC currently owns this IP on my segment?” — one address for the street, one for the world.

→ **Why that?** Why does TCP jitter? Its reliability machinery: a lost segment is retransmitted, and everything received after it **waits in a buffer** until the hole is filled (head-of-line blocking). A 1 kHz stream turns into bursts. Determinism and retransmission are structurally at odds.


→ **Why that?** Why do robot data streams pick UDP? Because a retransmitted sensor sample arrives *old*, and old is worse than absent — the same freshest-wins logic as the \$32 mailbox. ROS 2’s

DDS/RTPS rides on UDP and adds back exactly as much reliability as each topic asks for. For hard real-time over the wire, industry reshapes Ethernet itself: TSN time-slots it; EtherCAT slaves process the frame *as it flies through them*.

With numbers: a 1 kHz stream loses one packet; a 50 ms retransmission round-trip stalls the ~50 samples queued behind it — which then arrive as a burst the control loop never asked for. One loss, fifty late samples: head-of-line blocking, priced out.

	TCP	UDP
Model	Connection: 3-way handshake (SYN → SYN-ACK → ACK), ordered teardown	Connectionless datagrams — fire and forget
Guarantees	Ordered, complete, retransmitted; flow + congestion control	None: may drop, duplicate, reorder
Latency character	Bursty under loss (head-of-line blocking)	Steady; a loss costs one sample, not a stall
Overhead	20+ byte header, per-connection state on both ends	8-byte header, stateless
On a robot	Config, file/OTA transfer — anything that must arrive <i>complete</i>	Sensor/telemetry streams, video, DDS (ROS 2) — anything that must arrive <i>fresh</i>

Glue protocol	One-line job description
ARP	Resolves IP → MAC on the local segment (“who has 192.168.1.7?”)
DHCP	Leases the device an IP + gateway + DNS at boot — no lease, no network
DNS	Names → IPs; a device that hardcodes IPs breaks the day the server moves
NTP	Clock sync — the difference between mergeable logs and archaeology
ICMP	ping and unreachable messages — the first tool out when “it does not connect”



The three-way handshake — three segments, one agreement on sequence numbers.

Q: 1 kHz telemetry from a robot to a logger over flaky Wi-Fi. Transport?

- A. TCP — reliability is always better ❌ The word “reliable” doing the selling. Every loss stalls the stream to resend a sample that is stale on arrival.
- B. UDP — a lost sample is cheaper than a late stream ✅ Fresh beats complete for periodic data; add a sequence number if you want to count the losses.
- C. Raw Ethernet frames — skip the overhead ❌ Confined to one L2 segment, unroutable, and you reinvent addressing by hand. Real (EtherCAT lives here), but a specialist tool — not the default.
- D. HTTP — it is the standard ❌ Request/response for a push stream: headers dwarf the payload and the logger cannot be pushed to. Standard / suitable.

“How does TCP work?” — the 60-second spoken answer

TCP builds a reliable byte stream on top of unreliable IP. A connection opens with the three-way handshake — SYN, SYN-ACK, ACK — which synchronizes starting sequence numbers. From then on every byte is numbered: the receiver acknowledges what arrived, the sender retransmits what times out, and sequence numbers repair reordering. Flow control — the advertised receive window — keeps a fast sender from drowning a slow receiver, and congestion control backs the sender off when the network itself drops packets. The price of the guarantees is latency variance: one lost segment stalls delivery of everything behind it until the retransmission lands — which is exactly why freshness-critical streams choose UDP instead.

25 Wireless — the Range-Power-Bandwidth Triangle

Every radio question resolves to one triangle: **range, throughput, battery** — pick two. The physics behind it is energy per bit: slower symbols pour more energy into each bit, so the receiver can dig them out from further away (even from *below* the noise floor); fast symbols travel light and die young.

	BLE	Wi-Fi	LoRa / LoRaWAN	Zigbee / Thread
Range	~10–100 m	~50 m	2–15 km	~10–100 m, mesh-extended
Throughput	~100 kbps–1 Mbps usable	10s–100s of Mbps	0.3–50 kbps	~250 kbps
Battery	Coin cell, months–years	The battery assassin	Years on AA cells	Years (sleepy end devices)
Topology	Star (+ mesh profile)	Infrastructure AP	Star-of-stars via gateways	Self-healing mesh
Sweet spot	Phone-adjacent gadgets, config links, beacons	Bandwidth: video, OTA, robot backhaul	Sparse, far-flung sensors reporting rarely	Dense device networks: lighting, building automation

BLE in three sentences (the usual interview depth): devices **advertise**; a central connects to a peripheral and the link becomes a schedule of **connection events** — the connection interval is the power knob, since the radio sleeps between them. Data lives in a **GATT** server: services contain characteristics you read, write, or — the efficient path — subscribe to as **notifications**, so the peripheral pushes only on change.

LoRa’s party trick: chirp spread spectrum lets the receiver decode below the noise floor — that is where the kilometers-at-milliwatts come from. The bill: kilobits per second, regional duty-cycle limits (~1%), and downlink only in narrow post-uplink windows (Class A) unless you spend power listening (Class B/C).

Why? Why does Wi-Fi murder batteries even when “idle”? Staying associated means periodic beacons, key refreshes, and a power amplifier sized for megabits — the whole stack assumes wall power is nearby. BLE was designed from the opposite end: default radio-off, wake on a schedule.

Q: Soil-moisture sensor, one report per hour, 2 km from the nearest building. Radio?

- A. BLE **✗** *The low-power reflex — right instinct, wrong range by two orders of magnitude.*
- B. Wi-Fi with a high-gain antenna **✗** *Antennas do not repeal the link budget (the strict accounting: transmit power minus every loss must still exceed what the receiver can hear), and association upkeep alone would dominate an hourly-report power profile.*
- C. LoRaWAN **✓** *Kilometers at milliwatts for bytes per hour — the exact corner of the triangle LoRa was built to occupy.*
- D. Zigbee mesh **✗** *Mesh extends range through other nodes — which an empty field does not supply. Tempting because “mesh = range” hides the hop requirement.*

“How does BLE work?” — the 45-second spoken answer

A BLE peripheral advertises on three channels; a central scans, connects, and from then on both radios wake together only at scheduled connection events — the connection interval — and sleep in between, which is where coin-cell battery life comes from. Data is organized as a GATT server on the peripheral: services group characteristics, which the central reads, writes, or subscribes to as notifications, so values are pushed only when they change. Adaptive frequency hopping across 37 data channels rides out interference in the crowded 2.4 GHz band. Tens of meters of range, hundreds of kbps real-world.

26 Application Layer — MQTT, CoAP, HTTP & DDS

Everything below this layer moves bytes; this layer decides the **conversation pattern**. The embedded workhorse is MQTT: devices publish to **topics** on a central **broker**; consumers subscribe. Publisher and subscriber never know each other — decoupled in identity *and in time*, since the broker buffers across dropouts. A 2-byte minimum header, a **retained** message so late joiners instantly get the last value, and a **Last Will** the broker publishes for you when you die ungracefully — every design choice assumes flaky links and tiny devices.

Mapping example for one product: periodic telemetry → QoS 0 (a loss is fine, the next sample is coming); alarm events → QoS 1 with an idempotent handler; “unlock the door” commands → QoS 2 — because “at least once” on an actuator means it can happen twice.

QoS	Guarantee	Mechanism	Cost
0	At most once	Fire and forget	None — loss is possible
1	At least once	PUBLISH → PUBACK; retransmit until acked	Duplicates possible — the receiver must be idempotent
2	Exactly once	4-way: PUBLISH → PUBREC → PUBREL → PUBCOMP	Heaviest — the two extra legs exist purely to de-duplicate

	MQTT	CoAP	HTTP/REST	DDS (ROS 2)
Transport	TCP (+TLS)	UDP	TCP	UDP (RTPS)
Pattern	Pub/sub via broker	REST verbs for tiny devices, plus observe	Request / response	Brokerless pub/sub, automatic discovery
Can the server push?	Yes — the whole point	Yes (observe)	No — poll, or bolt on WebSocket	Yes, peer-to-peer
Sweet spot	Device fleets ↔ cloud	Constrained nodes, local REST	Config pages, OTA download, existing web APIs	On-robot: sensors, control, everything ROS 2

Why? Why pub/sub for machines instead of request/response? Producers and consumers scale independently — adding a tenth dashboard changes *nothing* on the sensor. And a device asleep behind NAT (the router trick that shares one public address among many devices — and blocks unsolicited inbound connections) cannot be polled, but it can always push when it wakes.

→ **Why that?** Why did ROS 2 choose DDS over an MQTT-style broker? A broker is a single point of failure and an extra hop — unacceptable in the control path of one robot. DDS discovers peers by multicast, and **every topic carries its own QoS**: reliability (best-effort vs reliable = the UDP/TCP trade, per topic), history depth (depth 1 = the §32 freshest-wins mailbox, now over the network), and deadline monitoring that flags a silent sensor. Concepts you learned per-bus reappear here as configuration.

Q: 1000 battery sensors reporting to the cloud over intermittent cellular. Protocol?

- A. HTTP polling from the server ❌ *Backwards twice: the server cannot reach sleeping NAT-ed devices, and headers dwarf 20-byte payloads.*
- B. MQTT, QoS 1 ✅ *Push on wake through the broker, 2-byte header, retransmit-until-acked with idempotent handling — designed for precisely this fleet.*
- C. CoAP ❌ *The respectable runner-up — genuinely built for constrained nodes — but UDP through carrier NATs is fragile and the broker/cloud ecosystem is thinner. Defensible, not the default.*
- D. Raw TCP sockets with a custom format ❌ *You would spend the project re-inventing QoS, wills, and reconnection — i.e. writing a worse MQTT.*

“How does MQTT work?” — the 60-second spoken answer

MQTT is publish/subscribe over TCP through a central broker. Each device holds one long-lived connection, publishes messages to hierarchical topics, and subscribes to the topics it cares about; the broker routes and buffers, so publishers and subscribers never know each other and need not be online at the same time. Three QoS levels trade overhead for guarantees: 0 is fire-and-forget, 1 retransmits until acknowledged — at least once, so handlers must be idempotent — and 2 runs a four-way exchange for exactly-once. A retained message hands late subscribers the latest value instantly, and the Last Will lets the broker announce a client that died without saying goodbye. Two-byte minimum header — built for tiny devices on flaky links.

PART VI • REAL TIME & RTOS

Who runs your code, which kernel to bet on, and determinism — the property that separates embedded from everything else.

27 Cores, Threads, Tasks & Processes — Who Actually Runs Your Code

Five words that interviews love to blur. Sort them once, by what each one actually owns:


Term	Hardware or software?	What it owns	On an MCU	On Linux
Core	Hardware	A complete CPU: registers, ALU, pipeline	Usually one; ESP32 has two (\$13)	Many — true parallelism scales with cores
HW thread (SMT)	Hardware	A duplicate register set sharing one core's execution units	Essentially absent	2 per core on hyperthreaded CPUs — looks like 2 cores, is not
Process	OS construct	A private address space + resources (files, permissions)	Does not exist — no MMU, one flat memory shared by all	The unit of isolation: one segfaults, the others live
Thread	OS construct	An execution context — register snapshot + its own stack — <i>inside</i> a process, sharing its memory	See task →	The unit the scheduler actually runs
Task (RTOS word)	RTOS construct	Register context + own stack + a priority	FreeRTOS “task” = Zephyr “thread”: a thread with no process around it — all share the flat memory	—

Term	Hardware or software?	What it owns	On an MCU	On Linux
ISR	Hardware-scheduled	Borrowed context on a borrowed stack (MSP)	Preempts everything; not schedulable, cannot block	Kernel-side interrupt context

ONE CORE — the scheduler interleaves: concurrency



TWO CORES — truly simultaneous: parallelism



One core interleaves — that is concurrency. Parallelism begins at core two.

Why? Why are there no processes on a bare MCU? A process is an **address-space illusion**, and the illusion requires an MMU to translate and police every access. Without one there is a single flat memory: any task can scribble on any other, so one bad pointer takes the whole system down — which is why MCU firmware leans so hard on watchdogs (§15) and code review. An MPU offers coarse region *protection*, but no translation — guard rails, not private worlds.

→ **Why that?** Why do “task” and “thread” mean the same thing on an RTOS? Because with no processes, a thread has nothing to live inside — it is just a register context, a stack and a priority. FreeRTOS called it a task, Zephyr calls it a thread; the concept is identical, and saying exactly that sentence dissolves the trick question.

→ **Why that?** Why is one core never parallel? One core executes one instruction stream; the scheduler interleaves contexts so fast it *feels* simultaneous — that is **concurrency**. **Parallelism** — two streams literally at once — needs a second core (or an SMT sibling, sharing the same execution units). All of your race conditions exist under mere concurrency; a second core just adds true simultaneity and, with it, the need for real memory-ordering care.

Q: FreeRTOS on a single-core Cortex-M4 runs two ready tasks. Are they executing at the same time?

A. Yes — that is what an RTOS is for **✗** *The marketing-level answer. An RTOS schedules; it cannot mint execution hardware.*

B. Yes — each task has its own stack, so they run independently **✗** *Own stack means own context, not own core. Independent state / simultaneous execution.*

C. No — one core, one instruction stream; the scheduler interleaves them (concurrency, not parallelism) **✓** *And the follow-up is free: parallelism starts at core number two — e.g. the ESP32 pinning radio to one core and the app to the other.*

D. Only if they share the same priority **✗** *Equal priorities enable time-slicing — which is still interleaving, just on a timer.*

“Process vs thread vs task?” — the 30-second spoken answer

A process is an address space plus resources — the unit of isolation, and it needs an MMU to exist. A thread is an execution context inside a process: its own registers and stack, sharing the process’s memory — the unit the scheduler runs. On a microcontroller there is no MMU, so there are no processes: everything is threads sharing one flat memory, and RTOS vendors just call them tasks. A context switch is therefore small — swap the register set and stack pointer — while a process switch on Linux also swaps the address space, which is why it costs more.

28 Embedded Linux vs RTOS

	RTOS (FreeRTOS, Zephyr)	Embedded Linux
Determinism	Hard guarantees — priority preemption, μ s-class, bounded latency	Soft — ms-class jitter; page faults (the MMU pausing a program mid-step to fetch memory it pretended was already there), scheduler fairness get in the way
Footprint	~10 kB RAM, boots in ms	~64+ MB RAM, MMU required, boots in seconds
You get	Predictability, tiny attack surface, runs on a \$1 MCU	Filesystems, full networking, Python/ROS, GPU drivers, security updates
Typical home	Motor control, sensor nodes, safety loops	Robot main computer, vision, gateways, HMI

Why? Why do real robots use **both**? Linux SBC runs ROS, planning and vision (feature-hungry, deadline-tolerant); an MCU with an RTOS runs the 10 kHz current loop and E-stop (deadline-critical, feature-poor). They talk over CAN/Ethernet/serial — describing this split answers "Linux or RTOS?" better than picking a side.

→ **Why that?** Why can't Linux just be made real-time? PREEMPT_RT gets it to ~tens of μ s worst case — good, but an MMU, caches and a fair scheduler mean the *worst case* is never as provable as a 200-line RTOS scheduler. Real-time is about proof, not averages.

29 The RTOS Landscape — FreeRTOS, Zephyr & the Others

Every kernel in this table schedules preemptive priority threads and offers queues, semaphores and mutexes — the §31 concepts map one-to-one across all of them. What you actually choose is the **ecosystem around the kernel**:

RTOS	Identity in one line	License / steward	Pick it when...
FreeRTOS	The minimal, ubiquitous kernel — you bring drivers and structure	MIT; stewarded by AWS	Smallest footprint, a port exists for everything, AWS IoT integration; the default "just a scheduler" choice
Zephyr	A full OS: kernel + device tree + driver model + BLE/net/FS subsystems + west build	Apache-2.0; Linux Foundation, multi-vendor	Product platforms that must outlive one chip: the device tree lets one app move across silicon vendors
Eclipse ThreadX (ex Azure RTOS)	Commercial-grade kernel with matched middleware (the reusable layers between kernel and app): FileX, NetX, USBX, GUIX	MIT; Eclipse Foundation (formerly Microsoft)	You want integrated, pre- safety-certified kernel + middleware artifacts (61508 / 26262 heritage)
μ C/OS-II / III	The textbook classic with medical/aero certification history	Open-sourced (was Micrium/SiLabs)	Legacy code bases, and its book remains the best kernel-internals teacher
NuttX	POSIX on a microcontroller — feels like a miniature Linux (/dev, standard APIs)	Apache-2.0; Apache Foundation	POSIX portability matters — notably the PX4 drone autopilot runs on it
RT-Thread	Kernel plus a huge package ecosystem and shell (FinSH)	Apache-2.0	Strong component marketplace; very large footprint in the Chinese ecosystem
VxWorks / QNX	The "big iron" commercial RTOSes: MMU processes, POSIX, microkernel (QNX)	Commercial	Safety-critical systems with budgets: Mars rovers flew VxWorks, cars ship QNX — these blur into §28's OS class

Why? Why so many kernels when the scheduler ideas are identical? Because the kernel is ~5% of the decision. The other 95% is drivers, middleware, build system, certification artifacts, and someone committed to maintaining it for your product's decade — you are choosing an **ecosystem and a governance model**, not a context-switch routine.

→ **Why that?** Why the Zephyr momentum? Vendor-neutral governance plus the device tree: hardware description moves out of C code, so the same application ports across vendors by editing a text file — the property Linux proved at scale, brought down to MCUs. The trade: a heavier build system (west, Kconfig) and a steeper first week than dropping three FreeRTOS files into a project.

→ **Why that?** How to say it in the room: “the APIs are isomorphic — task/queue/mutex map one-to-one (§32), so switching kernels costs drivers and build integration, not concepts. I would pick FreeRTOS for a small single-chip product, Zephyr for a multi-chip platform with a long life, ThreadX when certification artifacts must come with the middleware.” That sentence is the whole senior answer.

Q: One firmware platform must survive three different silicon vendors over a 10-year product line. Which RTOS bet?

A. FreeRTOS — it has the most ports **✗** *True fact, wrong conclusion: ports exist, but your drivers and board glue are still rewritten per chip — the porting cost lives outside the kernel.*

B. Zephyr — device tree + driver model abstract the silicon **✓** *Hardware description in data, drivers behind a stable API, vendor-neutral governance — built for exactly this bet.*

C. ThreadX — Microsoft heritage means longevity **✗** *Its actual differentiator is certified middleware, not portability — and its stewardship just changed hands, which cuts against the longevity story.*

D. Write a custom scheduler — full control **✗** *A weekend to write, a career to maintain: you would be re-deriving ten years of edge cases with a team of one.*

30 Deterministic Control-Loop Timing (the HAL_Delay Trap)


A junior implements a “500 Hz” PID like this:

```
while (1) {
    HAL_Delay(2);           /* blocks for >= 2 ms    */
    pid_update(&pressure_pid); /* takes t_exec      */
}
/* real period = 2 ms + t_exec + interrupt time -> < 500 Hz, and it wanders */
```

Why? Why can this never be 500 Hz? `HAL_Delay(2)` means “block for **at least** 2 ms”, and the delay only restarts *after* `pid_update` finishes. The period is 2 ms + execution time + whatever interrupts stole — always long, never constant. Error **accumulates**; there is no anchor to true time.

→ **Why that?** Why is jitter fatal to PID specifically? The D term is $\Delta\text{error}/\Delta t$ and the I term is $\text{error} \cdot \Delta t$ — the code assumes Δt is exactly 2 ms. When Δt breathes, your carefully tuned gains are effectively re-randomized every cycle: the loop hisses, oscillates, or goes limp.

→ **Why that?** Why does a hardware timer fix it? The timer counts in silicon, indifferent to what the code is doing. Its interrupt fires every 2.000 ms regardless of execution paths — the schedule is anchored to the crystal, not to how long last cycle took. Determinism holds as long as the loop’s WCET (worst-case execution time) stays under 2 ms.



Only one of these loops is anchored to the crystal — the other accumulates its own execution time.

Correct shapes — bare-metal and RTOS:

```

/* bare-metal: hardware timer TIM6 configured for 500 Hz */
volatile uint8_t tick;
void TIM6_IRQHandler(void) { tick = 1; }

while (1) {
    if (tick) { tick = 0; pid_update(&pressure_pid); } /* anchored to the crystal */
}

```

```

/* FreeRTOS – absolute deadlines */
TickType_t t = xTaskGetTickCount();
for (;;) {
    vTaskDelayUntil(&t, pdMS_TO_TICKS(2)); /* wakes AT t+2, t+4 ... never AFTER work */
    pid_update(&pressure_pid);
}

/* Zephyr – periodic kernel timer, same guarantee */
K_TIMER_DEFINE(loop_timer, NULL, NULL);
k_timer_start(&loop_timer, K_MSEC(2), K_MSEC(2));
for (;;) {
    k_timer_status_sync(&loop_timer); /* blocks until the next 2 ms expiry */
    pid_update(&pressure_pid);
}


```

Note `vTaskDelayUntil` vs `vTaskDelay`: *until* targets an absolute next wake time (period stays 2 ms and error cannot accumulate); plain *delay* is HAL_Delay’s mistake wearing an RTOS badge — and Zephyr’s `k_sleep(K_MSEC(2))` is the same trap in Zephyr clothing. The drift-free Zephyr forms: the periodic `k_timer` above, or an absolute timeout via `K_TIMEOUT_ABS_TICKS`. The gold-plated answer: timer-triggered ADC via DMA, PID in the conversion-complete callback — sampling itself becomes jitter-free.

31 The Synchronization Toolbox — Mutexes, Semaphores & Friends, In Depth

Preemption (§27) means two tasks can be interleaved at *any* instruction; shared data means §6’s torn values are waiting to happen. Synchronization primitives are the kernel’s answers — each one a differently shaped tool. But first, the picture every one of them stands on: the life of a task.

The task lifecycle — where “waiting” actually goes



Four states, five arrows — every primitive below is just a different reason to take the red arrow and a different doorbell for the green one.

Why? Why is **blocking** the whole point of an RTOS? A BLOCKED task is parked: it costs **zero CPU** while it waits, and the scheduler spends that time on whoever is READY. The alternative — spinning on a flag in a loop (§39) — burns the entire processor to discover “not yet”. Also note the subtle arrow: waking lands you in READY, not RUNNING — the scheduler still gets the final word.

Semaphores — the doorbell and the parking garage

A **binary semaphore** is a doorbell with a one-ring memory. `give()` rings it: an internal flag is set and the first parked waiter is moved to READY. `take()` answers it — or, if nobody has rung yet, parks the caller BLOCKED until someone does. No data travels; *the ring itself is the message*. And because nobody “owns” a doorbell, **anyone may give it — including an ISR**, which is exactly its §32 job of waking a task from an interrupt.

A **counting semaphore** is the parking garage: the count is the number of free spaces. `take()` drives in (count decreases, or the caller waits at the barrier when it reads 0); `give()` drives out (count increases). Two distinct uses: guarding **N identical resources** (four DMA channels → create it with count 4), and **counting events so bursts survive** — five ISR gives while the task is busy leave count = 5, and the task then drains all five. A binary semaphore would have coalesced them into a single ring: five events, one wake, four silently lost. That coalescing is the exact trap §32’s event-group question turns on.

Mutex — the key with a name written on it


A mutex looks like a binary semaphore until you notice one added word: an **owner**. `lock()` writes the caller’s name on the key; only the owner may unlock it; most kernels let the owner re-lock recursively, counting the nesting. Ownership is not bureaucracy — it is the enabling fact behind the cure in the next diagram.

Why? Why can a semaphore never offer priority inheritance? Inheritance means “boost **the holder** to the waiter’s priority” — and a semaphore has no holder. Anyone may give it, so the kernel literally does not know whose priority to raise. No owner, no cure.


→ **Why that?** Why is a mutex banned inside an ISR? Both halves fail: an ISR cannot **block** (it has no task state to park — §27), and it cannot **own** (the owner field expects a task to boost). §32’s rule — semaphore or notification from ISRs, mutex between tasks — falls straight out of the definitions.

Priority inversion — the failure, then the cure

WITHOUT inheritance (binary semaphore guarding the resource)



WITH inheritance (mutex): the holder borrows the waiter’s urgency




= takes lock = releases hatched = blocked waiting for the lock time →

The same three tasks, the same lock — ownership plus inheritance turns an unbounded stall into a short, computable one.

The top panel is the sequence that rebooted the Mars Pathfinder (§39): H ends up waiting on M, a task that holds nothing and owes nobody. With a mutex, the instant H blocks the kernel boosts L to H’s priority — M cannot butt in, L finishes its short critical section, and H’s worst-case stall becomes a computable number. The stricter cousin, **priority ceiling**, boosts the holder immediately to the lock’s highest possible user — inversion prevented rather than cured.

The rest of the toolbox, briefly — a queue is the conveyor: it carries the data and the wake-up in one operation (its field manual is §32). An event group is a checklist of bits a task can wait on — ANY or ALL — the clean answer to “proceed when pump AND motor AND sensors are ready”. And the length-1 overwrite mailbox keeps only the freshest value — the pattern §24 and §32 keep returning to. When in doubt, walk the flowchart:



One question at a time: data → thing → conditions → counting — and the plain doorbell if all answers were no.

Q: Two tasks share an I2C bus guarded by a binary semaphore. Occasionally the highest-priority task stalls for many milliseconds waiting for the bus — while an unrelated medium task runs. Root cause and fix?

- A. The bus is slow — raise I2C to 400 kHz **✗** Speed shortens the hold time, not the scheduling hole — and the stall tracking the MEDIUM task, not the bus, is the tell-tale.
- B. Priority inversion: the low-priority holder got preempted by medium, and a semaphore cannot boost it. Replace the guard with a mutex — priority inheritance **✓** Names the mechanism (the top panel above) and why the fix works: ownership makes the boost possible.
- C. Raise the high task's priority further **✗** It is already the highest — inversion is precisely the situation where being highest does not help you.
- D. Have the I2C ISR give the semaphore to free the bus sooner **✗** A give is not a release-of-claim: a stranger's give corrupts the guard while the real user is mid-transaction — the no-owner weakness dressed up as a favor.

"Mutex vs semaphore?" — the 45-second spoken answer

Both can park a task; they answer different questions. A semaphore is a **signal**: a doorbell (binary) or a counter of events and free resources (counting). Anyone may give it — which is exactly why it is the tool for ISR-to-task wake-ups — and it has no concept of who is holding anything. A mutex is **exclusive ownership** of a shared thing: the locker's name is recorded, only the owner may unlock, and that ownership is what makes priority inheritance possible — the holder is temporarily boosted so a high-priority waiter can never sit behind an unrelated medium task. So: to signal, semaphore; to protect, mutex; and never a mutex from an ISR, which can neither block nor own.

32 RTOS in Practice: Losing Data & Choosing Primitives (FreeRTOS / Zephyr)

Scenario: sensor task pushes a reading to a 10-deep queue every 1 ms; a lower-priority PID task consumes. Under an interrupt storm, readings vanish silently and the PID goes erratic.

Why? Why do readings vanish with no error? The consumer is starved, the queue fills in 10 ms, and every further `xQueueSend(q, &v, 0)` returns `errQUEUE_FULL` — a return value nobody checked. FreeRTOS did exactly what it was told; "silent" loss is an **unchecked return code**, not a kernel bug. Zephyr twin: `k_msgq_put(&q, &v, K_NO_WAIT)` returns `-ENOMSG` when full — the same silence if unchecked.

→ **Why that?** Why is "make the queue deeper" only a bandage? Depth buys time, not correctness — a longer storm still overflows it, and now the PID acts on readings 40 ms stale. For control, old data is worse than lost data.

→ **Why that?** Why is the strong answer a depth-1 mailbox? A control loop wants the **fresh**est sample, not a backlog: `xQueueOverwrite()` on a length-1 queue makes the newest value always win, loss becomes explicit policy instead of accident — plus fix the starvation itself (priorities, shorter critical sections). Zephyr's `k_msgq` has no overwrite; the idiomatic freshest-value answer there is a **zbus**

channel — latest-message semantics by design — or purge-then-put.

Q: Best explanation of the failure and fix?

A. Queues are thread-safe, so loss is impossible — it must be float precision **✗** *Thread-safe / loss-proof: safety governs concurrent access, not capacity. The red herring works because the first clause is true.*

B. Starved consumer → queue full → xQueueSend with 0 timeout returns errQUEUE_FULL, unchecked; fix priorities / handle the return / overwrite-oldest **✓** *Names the mechanism, the silent part (unchecked return), and layered fixes. Zephyr equivalent: k_msgq_put → -ENOMSG.*

C. FreeRTOS queues are kernel-limited to 10 items; replace with a custom ring buffer **✗** *A confident fake constraint — depth is whatever you pass to xQueueCreate. Beware invented "kernel limits" stated with authority.*

D. Use a volatile global instead — queues are too slow for 1 ms data **✗** *The §6 lesson weaponized: volatile gives no atomicity, no ordering, no wake-up. This "fix" trades visible loss for invisible races.*

Choosing the right primitive — the four situations that cover 90% of designs:


You need...	Reach for — FreeRTOS / Zephyr	Why — and the trap
ISR wakes a task ("CAN frame arrived")	Binary semaphore — <code>xSemaphoreGiveFromISR()</code> / <code>k_sem_give()</code> . Lighter still: task notification / <code>k_event</code>	Zephyr has no FromISR twins — the same <code>k_sem</code> API is ISR-callable (with <code>K_NO_WAIT</code>). Never a mutex from an ISR in either kernel: mutexes have owners — a meaningless concept for an ISR
Pass data, in order, one at a time	Queue — <code>xQueueSend/Receive</code> / <code>k_msgq</code> (copies) or <code>k_fifo</code> (pointers)	Carries the payload <i>and</i> the synchronization in one operation. A semaphore can only say "something happened", never <i>what</i>
Two tasks share one SPI bus / I2C bus / buffer	Mutex — <code>xSemaphoreCreateMutex()</code> / <code>k_mutex</code>	Ownership + priority inheritance in both kernels — the holder is temporarily boosted so a high-priority waiter isn't stuck behind a medium task (§39 priority inversion). A binary semaphore lacks exactly this
Wait until pump AND motor AND sensors are all initialized	Event group — <code>xEventGroupWaitBits()</code> / <code>k_event_wait_all()</code>	Each subsystem sets its bit; the supervisor blocks until all three are set. One binary semaphore given 3× cannot count — gives coalesce, the supervisor releases early. That subtlety is the whole question
Consumer only ever needs the latest value	Length-1 queue + <code>xQueueOverwrite</code> / a <code>zbus</code> channel	The "mailbox" pattern from the scenario above — freshest-wins by design (zbus keeps only the latest message per channel)

Why? Why does Zephyr skip the FromISR split? Its kernel checks the calling context at runtime — one API surface, with blocking simply forbidden (`K_NO_WAIT`) inside an ISR. FreeRTOS splits the API instead, so the task-context fast path never pays for that check. Explaining *that* trade-off is the cleanest "I have shipped on both kernels" signal you can send.

The difference in four lines — giving a semaphore from an ISR in both kernels:

```
/* FreeRTOS — the FromISR ceremony */
 BaseType_t woken = pdFALSE;
 xSemaphoreGiveFromISR(rx_sem, &woken);
 portYIELD_FROM_ISR(woken);      /* context-switch NOW if a higher task woke */

/* Zephyr — the same API as thread context */
 k_sem_give(&rx_sem);             /* ISR-safe as-is; the kernel checks context */
```



a "context" is exactly these registers + this stack — nothing more

A context switch is one pointer swap with ceremony: save PSP, load PSP, pop.

"How does an RTOS context switch work?" — the 60-second spoken answer (Cortex-M)


A tick or an ISR decides a higher-priority task is ready and pends PendSV — an exception deliberately set to the **lowest** priority, so the switch only happens after every real interrupt has finished. On exception entry the hardware has already stacked R0-R3, R12, LR, PC and xPSR onto the running task's process stack; the PendSV handler pushes the remaining R4-R11, saves that stack pointer into the task's control block, loads the next task's saved stack pointer, pops its R4-R11 and returns — hardware unstacks the rest and the new task resumes mid-thought. A context is nothing more than a register set plus a stack; the scheduler is bookkeeping around one pointer swap. That is also why §27's task switch is cheap while a Linux process switch — which swaps address spaces — is not.

PART VII • THE ROBOTICS LAYER

Control, sensing and motion — the layer that turns firmware into a robot.

33 PID Control — the Question Robotics Interviews Always Reach

"Implement a PID controller" and "how would you tune it" both appear verbatim in interview reports. The controller in one line: **error = setpoint - measurement**, and three terms each look at a different tense of that error.



The PID loop: three views of one error, summed into the actuator, closed by the sensor.

Term	Looks at	What it fixes	Too much of it causes
P ($K_p \cdot e$)	The error now	Responsiveness — the muscle	Overshoot, oscillation; P alone always leaves a steady-state offset under load
I ($K_i \cdot \Sigma e \cdot \Delta t$)	The accumulated past	Erases steady-state offset — keeps pushing while any error persists	Windup : the sum grows huge while the actuator saturates, then blows through the setpoint
D ($K_d \cdot \Delta e / \Delta t$)	The trend — where error is heading	Damping: brakes before arrival, calms P's oscillation	Amplifies sensor noise — derivative of noise is chaos; always filter it

```

e      = setpoint - y;
integ += e * DT;
if (integ > I_MAX) integ = I_MAX;      /* anti-windup: clamp the accumulator */
if (integ < -I_MAX) integ = -I_MAX;
out    = KP*e + KI*integ + KD*(e - e_prev)/DT;
e_prev = e;                          /* DT must be CONSTANT — which is §30's */
                                         /* deterministic-loop story, now load-bearing */

```

Why? Why does P alone settle below the setpoint? At the setpoint the error is zero, so P's output is zero — but holding a load (gravity on an arm, friction on a wheel) needs a *nonzero* output. The system slides down until the leftover error, times Kp, exactly pays the holding cost. Only a term that remembers — I — can output effort at zero present error.

→ **Why that?** Why is windup the classic field failure? During a big move the actuator saturates; error persists; the integral quietly accumulates a mountain. When the target is finally reached, that mountain must drain — as a massive overshoot. Fix: clamp the integrator, or stop integrating while saturated. Naming windup unprompted is a strong-hire tell.

→ **Why that?** Why does the tuning ritual go P → D → I? Raise Kp until it responds briskly and just starts to ring; add Kd to damp the ring; add the smallest Ki that removes the leftover offset. Each term is tuned against the misbehavior the previous one introduced.

Q: A wheel-speed loop responds quickly but settles 5% below the setpoint whenever the robot carries a load. First fix?

A. Increase Kp **✗** *The reflex — and it does shrink the offset (offset = load/Kp) — but never to zero, and enough Kp to make it negligible buys you oscillation instead.*

B. Add / increase Ki **✓** *A persistent offset under load is the signature of a missing integral term: only accumulated error can fund a standing output at zero error.*

C. Add Kd **✗** *D shapes transients; it is blind to a constant offset — the derivative of a constant is zero.*

D. Double the sample rate **✗** *Faster sampling refines the same wrong answer. Timing matters for stability (§30), not for steady-state offset.*

“How does a PID controller work?” — the 45-second spoken answer

Every cycle I compute the error — setpoint minus measurement — and sum three views of it. Proportional pushes against the error right now: the muscle. Integral accumulates past error, so it can hold a standing output once the error is gone — that is what removes steady-state offset under load, and it needs anti-windup clamping for when the actuator saturates. Derivative looks at the error's slope and brakes before arrival — damping, but it amplifies noise, so it gets filtered. The sum runs at a fixed, deterministic Δt — the gains are only meaningful if the loop period is — and tuning goes P for response, D to calm it, then the least I that kills the offset.

34 Motors & Encoders — Making and Measuring Motion


	Brushed DC	BLDC	Stepper	Servo (integrated)
Commutation	Mechanical brushes	Electronic — the controller must know rotor position (Hall sensors, encoder, or back-EMF — the voltage the spinning motor itself generates: free position information)	None — you step the field, the rotor follows	Whatever is inside + its own controller
Feedback	None by default	Position sense required to commutate; encoder for control	Open loop — assumed, not measured	Built in — closed loop out of the box

	Brushed DC	BLDC	Stepper	Servo (integrated)
Character	Cheap, simple, brushes wear and spark	Efficient, long-lived, power-dense — needs an ESC/FOC (§13's motor timers)	Exact positioning at low speed, torque collapses with speed	A subsystem, not a motor: you command, it executes
Sweet spot	Toys, pumps, anything cost-floor	Drones, drive wheels, spindles — modern robotics default	Printers, small axes — where steps are cheap and loads known	Robot arms/joints where integration beats cost

Why? Why do steppers “lose steps” — the reported interview probe? Open loop means the controller *assumes* every commanded step happened. Exceed the torque the magnetic field can deliver — a snag, too much acceleration — and the rotor stalls a step behind while the controller's position count marches on. Nothing detects it; the error is permanent until re-homing. Closed-loop steppers exist precisely to close this hole.

→ **Why that?** Why does BLDC efficiency cost you a whole control problem? Removing brushes removes the component that *mechanically* knew when to switch the coils — so the electronics must recreate that knowledge from rotor position, which is the entire job of an ESC and, done properly with sinusoidal currents, of FOC (§13's PWM-synchronized ADC exists for this).

Encoders — how the robot knows where it is: an **incremental** encoder outputs two square waves, A and B, 90° apart (quadrature). Which one leads gives **direction**; counting every edge of both channels gives **4×** the line count in resolution; a once-per-rev index (Z) pulse gives a homing reference — because incremental encoders forget everything at power-off. An **absolute** encoder reports true position from the first read — its discs are patterned in **Gray code**, where exactly one bit changes per step, so a read taken mid-transition can never be off by more than one count (a binary-coded disc could momentarily read wildly wrong when several bits flip together).



Quadrature: the 90° phase encodes direction; four edges per line give the $\times 4$ resolution.

Q: A 1024-line incremental encoder is decoded in full quadrature ($\times 4$). Counts per revolution?

- A. 1024 **✗** Reads the datasheet line count and stops — ignoring that A and B each contribute two edges per line.
- B. 2048 **✗** Half-quadrature (counting one channel's both edges) — a real mode, but not what $\times 4$ means.
- C. 4096 **✓** Two channels $\times$ two edges per line: 4 countable events per line, $1024 \times 4 = 4096$. On STM32, the timer's hardware encoder mode does this counting for free.
- D. 8192 **✗** Over-multiplication — there are only four state transitions per line in a quadrature cycle, not eight.

Odometry example: 100 mm wheels, 4096 counts/rev $\rightarrow \pi \cdot 100 / 4096 \approx 0.077$ mm per count — lovely resolution, and yet 1% slip over a 10 m run is 100 mm of confidently wrong position. That single number is §35's entire case for fusion.


35 Sensor Fusion — Why Robots Never Trust One Sensor

The reported testimonial stack — wheel odometry + IMU + lidar/camera fused in an EKF — exists because every sensor lies in its own characteristic way:

Sensor	What it measures	Its virtue	Its lie
Gyroscope	Angular rate	Smooth, fast, immune to vibration direction	Integrating rate to get angle also integrates its bias (the small constant error a real gyro reports even sitting perfectly still) — the estimate drifts without bound
Accelerometer	Specific force — gravity's direction when still	Absolute tilt reference: drift-free	Noisy, and any real acceleration masquerades as gravity — useless mid-maneuver
Wheel encoders (odometry)	Wheel rotation → distance	Precise over short runs, high rate	Slip and skid accumulate — confidently wrong after enough meters
Magnetometer	Heading vs magnetic north	Absolute yaw — the axis gravity cannot give	Every motor and steel beam bends its truth

Why? Why does fusion work when each sensor alone fails? The lies are **complementary in frequency**: the gyro is right short-term and wrong long-term; the accelerometer is the exact opposite. Trust the gyro for fast changes and let the accelerometer slowly correct the drift, and both weaknesses cancel. That single sentence is the complementary filter:

```
angle = A * (angle + gyro_rate * dt)    /* high-pass: trust gyro for the fast part */
      + (1 - A) * accel_angle;          /* low-pass: let gravity fix the drift    */
/* A ~ 0.98 : the crossover — one tunable, runs on an 8-bit MCU */
```



Complementary filter as blocks: gyro carries the fast path, gravity slowly re-anchors it.

Why? Why then a Kalman filter? The complementary filter's A is a *fixed* trust ratio. A Kalman filter makes trust **dynamic**: it tracks an uncertainty for its own estimate, grows it during prediction, and on each measurement computes the optimal blend (the Kalman gain) from the two uncertainties — trusting the sensor more when the estimate is shaky and vice versa. The **EKF** is the same loop with the nonlinear robot model linearized around the current estimate each step — which is why every localization stack in the testimonials says “EKF”.

Q: Why can't a gyroscope alone give you tilt angle, if you can just integrate its rate?

- A. Gyros are too noisy ✗ Backwards — gyros are the quiet sensor; noise is the accelerometer's vice. The confusion of the two is the trap.
- B. Integration also accumulates the bias, so the angle drifts without bound ✓ A constant 0.1°/s bias is 6° of error per minute, forever growing. Drift, not noise, is the gyro's failure mode.
- C. Rate sensors cannot be integrated in fixed-point math ✗ Technical-sounding filler — integration is one multiply-accumulate.
- D. The gyro only works while the robot moves ✗ Gyros read zero rate at rest just fine; it is the angle bookkeeping that decays.


“How does a Kalman filter work?” — the 45-second spoken answer

It is a two-step loop. Predict: push the state forward with your motion model — and grow its uncertainty, because prediction is guesswork. Correct: take a measurement, which has its own uncertainty, and blend the two — the Kalman gain is computed from the ratio of the uncertainties, so a confident estimate barely moves for a noisy sensor, and a shaky estimate snaps toward a good one. Repeat every cycle. For nonlinear systems like robots, the EKF linearizes the model around the current estimate at each step. In one sentence: it is the complementary filter’s trust knob, made optimal and self-adjusting.

36 Kinematics & the ROS Toolbox

Forward kinematics (FK): given the joint angles, where is the end effector? Chain the frame transforms link by link — one unique answer, pure bookkeeping (the DH convention is just a standard way to parameterize those transforms). **Inverse kinematics (IK):** given a desired pose, what joint angles achieve it? Now the equations are nonlinear and run backwards — there may be **several** solutions (elbow-up/elbow-down), **none** (out of reach), or a continuum (redundant arms) — which is why IK is solved numerically or with per-robot analytic derivations, and why a general 6-DOF pose needs six joints.

Worked FK check (unit links): $\theta_1 = 90^\circ$, $\theta_2 = -90^\circ \rightarrow$ elbow at (0, 1), end effector at (1, 1). Have one of these memorized — thirty seconds of cos/sin on the whiteboard that proves the transform chain actually works.



FK: $(\theta_1, \theta_2) \rightarrow$ exactly one pose. Bookkeeping.

IK: pose $\rightarrow (\theta_1, \theta_2) \dots$ solid red = elbow-up, dashed gray = elbow-down — same target, two answers. Move the target out of reach: zero answers. That structure is why IK is hard.

FK maps angles to one pose; IK runs backwards into elbow-up / elbow-down.

Why? Why does the **Jacobian** keep appearing? It is the local bridge between the two worlds: it maps joint velocities to end-effector velocity. Where it loses rank (rank = how many independent directions of motion it can still produce) — a **singularity** — the arm has locally lost a direction of motion, and near it the inverse demands absurd joint speeds for tiny end-effector moves. “What is a singularity?” = “the Jacobian dropped rank; a straight wrist is the classic example.”

→ **Why that?** Why is everything in robotics “frames and transforms”? Because every sensor and actuator reports in its own coordinate frame, and fusing them (§35) requires expressing all of it in one. ROS makes this literal: **tf2** maintains the tree of time-stamped transforms between frames, and asking “where was the camera relative to the base 80 ms ago” is a library call, not a derivation.

The ROS 2 vocabulary interviews expect, one line each: **nodes** are processes; **topics** stream (pub/sub over $\$26$ ’s DDS), **services** request/reply, **actions** run long goals with feedback; **URDF** describes the robot’s links, joints and inertias (what simulators like Gazebo load); **launch files** compose the node graph; **bags** record topics for replay — the debugging superpower. Path-planning one-liners for the same interviews: **A*** — optimal grid search with an admissible heuristic; **RRT** — random sampling that scales to high-DOF arm spaces; **SLAM** — the chicken-and-egg of building the map while localizing inside it, attacked with EKFs or pose-graph optimization.


Q: Why is inverse kinematics fundamentally harder than forward kinematics?

- A. It involves matrix inversion, which is expensive **✗** *The name-bait: only the velocity-level problem inverts the Jacobian — and even that fails exactly at singularities. Position IK is not “FK, inverted”.*
- B. The equations are nonlinear with multiple, one, or zero solutions **✓** *FK is a unique chained transform; IK runs the nonlinearity backwards — elbow-up vs elbow-down, unreachable poses, redundant continua. Structure, not arithmetic, is the difficulty.*
- C. FK requires sensors and IK does not **✗** *Neither requires sensors — both are geometry. Encoder confusion imported from §34.*
- D. IK is only hard for more than 6 DOF **✗** *Redundancy (>6) adds infinite solutions, but a plain 6-DOF arm already has the multiple/zero-solution structure.*

PART VIII · SENIOR-LEVEL JUDGMENT

System trade-offs, fault reasoning, and the rapid-fire traps.

37 Design Factors in Embedded Systems




Every ring must hold: a fast MCU (core) is worthless if the business ring (cost, lifecycle) breaks.

Layer	You must decide	The "why" an interviewer wants
Hardware	MCU, memory sizes, power budget, I/O count	Everything above inherits these ceilings — you cannot software your way past 2 kB of RAM
Firmware	RTOS vs super-loop, interrupt strategy, secure boot	Concurrency model determines every timing bug you will ever debug
System	Reliability targets, watchdogs, OTA update, compliance (CE, ISO 26262)	Field failures cost 100× more than design-time fixes
Performance	WCET (worst-case execution time), throughput vs latency	Real-time means the worst case meets the deadline — averages are irrelevant to the airbag
Business	BOM cost (bill of materials — the parts list), manufacturability, 10-year part availability	The best design that misses its price point ships zero units

Why? Why "throughput vs latency" as a named trade-off? Buffering/batching raises throughput but adds delay. A logger wants throughput; a current-control loop at 20 kHz wants latency. Same silicon, opposite tuning.

38 Sensor Reasoning: Diagnosing Faults from Two Signals

A favorite written question: *pump* → *hose (flow sensor)* → *pressure sensor* → *nozzle*. Detect a nozzle clog, and distinguish it from an empty paint tank — both stop paint from spraying.



empty tank instead: the pump churns air — pressure AND flow collapse together. Two sensors, unambiguous verdict.
 positive-displacement pump keeps forcing volume into a blocked line: pressure climbs while flow dies — the divergence.

The divergence signature: pressure and flow split only when the path is blocked.

Condition	Pressure (pre-nozzle)	Flow (hose)	The physics
Normal spray	Nominal	Nominal	Steady flow through nozzle restriction sets a steady back-pressure
Nozzle clog	Rises sharply	Falls to ~0	A peristaltic pump is positive-displacement: it keeps forcing volume into a dead-ended line, so pressure climbs while nothing moves
Tank empty	Falls / erratic	Falls to ~0 (or noisy air readings)	Nothing to pressurize: the pump churns air, so both signals collapse together

Why? Why do the two sensors *diverge* only for the clog? Because a blockage separates "energy applied" (pressure) from "material delivered" (flow). Starvation kills both at once. The diagnostic is the **pattern across sensors**, never one threshold.

→ **Why that?** Why do interviewers love this shape of question? It is the essence of embedded fault detection: cheap sensors + physics reasoning > one expensive sensor. The same divergence logic finds blocked air filters (ΔP high, flow low) and dry-running water pumps (current low, flow zero).

Model answer skeleton (3–5 sentences)

Name the pump type and its consequence (positive displacement → constant-flow source). State each sensor's reading *and the mechanism*: clog → pressure integrates upward against a sealed line while flow drops to zero. Give the discriminator: clog = pressure UP + flow DOWN; empty tank = BOTH down. Close with the action: threshold on the divergence, alarm, stop the pump before hose/pump damage.

39 Rapid-Fire Round — Frequently Sprung Traps

Topic	The 20-second answer that lands
Endianness	Little-endian stores the least-significant byte at the lowest address (ARM default, x86); big-endian the opposite (network byte order). Why asked: the moment two systems exchange multi-byte data (CAN payloads, TCP), a swapped uint32 becomes garbage — hence htonl()/ntohl().
Mutex vs semaphore	Mutex = lock with an owner , for mutual exclusion; good ones add priority inheritance. Semaphore = counter for signaling (ISR → task) or counting resources. Trap: "a mutex is a binary semaphore" — no: ownership and priority inheritance are the difference. APIs: FreeRTOS xSemaphoreCreateBinary / ...CreateMutex · Zephyr k_sem / k_mutex.
Priority inversion	Low-priority task holds a mutex the high-priority task needs; a medium task preempts the low one, so high waits on medium. Killed the Mars Pathfinder until NASA enabled priority inheritance (holder temporarily runs at waiter's priority). Guaranteed follow-up to any mutex answer.
ISR rules	Keep it short; no blocking calls, no malloc, no printf; set a flag / push to a queue and let a task do the work ("defer to bottom half"). Shared data → volatile + critical section. Why: while you're in an ISR of priority N, everything ≤ N waits — you are stealing time from the whole system. API note: FreeRTOS allows only *FromISR calls in an ISR; Zephyr uses the same APIs everywhere but only with K_NO_WAIT.

Topic	The 20-second answer that lands
Interrupt latency	Time from event to first ISR instruction: hardware sync + current-instruction completion + context save + vector fetch (~12 cycles on Cortex-M) — plus any time interrupts were masked, which your critical sections control. Longest critical section bounds your worst case.
DMA	Peripheral ↔ memory copies without the CPU: ADC streams to a buffer while the CPU computes. Why asked: shows you can move beyond "CPU does everything" — plus the classic trap that the CPU cache may not see DMA-written RAM (cache coherency).
static (3 meanings)	(1) inside a function: value persists across calls; (2) global/function scope: hidden from other files (linkage control); (3) in C++ classes: shared across instances. Interviewers ask for all three; most candidates give one.
Reentrancy	A function safe to re-enter mid-execution (from ISR or another task): no static/global state, no non-reentrant calls. Trap answer: strtok() is the canonical non-reentrant offender.
Pull-up / pull-down	A floating input pin picks up noise and reads randomly; a resistor defines the idle level. Also the enabler of open-drain buses (I2C, one-wire) and shared interrupt lines.
Cross-compilation	Compile on x86 host, run on ARM target — the target can't host a compiler. Toolchain triplet (arm-none-eabi-gcc), linker script places .text/.data into the MCU memory map you saw in §7.
Zephyr vs FreeRTOS (topical)	FreeRTOS = minimal kernel, you bring drivers. Zephyr = kernel + device tree + driver model + build system (west) — a "Linux-like" workflow for MCUs. Why asked: shows you follow where the industry (and ROS 2 micro-controllers via micro-ROS) is heading.
ROS pub/sub vs services	Topics: many-to-many streaming, fire-and-forget (sensor data). Services: request/reply, blocking (trigger a calibration). Actions: long-running with feedback (move arm to pose). Picking the wrong one is the reviewer's favorite design question.

Round 2 — harvested from interview write-ups

Topic	The 20-second answer that lands
Microcontroller vs microprocessor	MCU = CPU + flash + RAM + peripherals on one die, runs standalone (a system on a chip in the small sense); MPU = the CPU alone, needs external memory and support — and usually an MMU, so it runs Linux-class software. "Which did you pick and why" is the real question behind it.
Deadlock	Two tasks each hold a lock the other needs — forever. All four Coffman conditions must hold: mutual exclusion, hold-and-wait, no preemption, circular wait. Practical cure: break the circle — acquire locks in one global order, or take them with timeouts.
Semaphore vs spinlock	Both guard resources; a blocked semaphore sleeps the task (scheduler runs someone else), a spinlock burns CPU polling. Spinlocks only make sense on multi-core for critically short sections — on a single-core MCU, spinning on a lock a same-core task holds is a deadlock with extra steps.
MMU vs MPU	MMU translates addresses → private address spaces → processes and virtual memory. MPU only <i>checks</i> accesses against a few regions — protection without translation. That one difference is why Linux needs an MMU and §27's MCU has no processes.
Switch debouncing	Contacts physically bounce for ~1–10 ms, so one press reads as many edges. Fix in hardware (RC + Schmitt) or software: on edge, start a ~20 ms timer and accept the level that survives it. "Interrupt per edge + timer confirm" is the answer they want.
Nyquist & aliasing	Sample faster than 2× the highest frequency <i>present</i> — not just the one you care about — or high frequencies fold down and impersonate low ones. That is why an anti-alias low-pass sits before the ADC: after sampling, the damage is indistinguishable from signal.
ADC resolution vs accuracy	LSB = $V_{ref} / 2^n$: a 12-bit ADC at 3.3 V resolves ~0.8 mV steps. Resolution is step size; accuracy is how true those steps are (noise, INL/DNL, reference drift) — a 16-bit ADC with a sloppy reference is a precise liar. Recall §13: this is the ESP32's ADC complaint.

Topic	The 20-second answer that lands
Safe string functions	strlen/strcpy trust a terminator that hostile or corrupted data may not have — unbounded reads/overflows. Use the bounded forms (strlen, strcpy/sprintf) and say <i>why</i> : every buffer gets an explicit size the function cannot exceed.
Recursion in embedded	Legal but shunned: stack depth becomes input-dependent, so WCET and stack sizing (§7) become unprovable. MISRA bans it. Convert to iteration or an explicit stack with a known bound.
malloc policy on MCUs	If used at all: allocate everything at init, never in the steady state — or replace with static pools / fixed-size block allocators. The reason is §7's fragmentation story; saying "pools" marks field experience.
Is i++ atomic?	No — it is load, add, store; an ISR between them loses an update even with volatile (§6). Protect with a critical section or use hardware atomics (LDREX/STREX on ARM).
Safe OTA update	Two slots so the only good image is never erased; verify (CRC/signature) before marking bootable; watchdog + rollback flag demotes an image that boots but hangs. The design goal has a name: anti-bricking (§12).
JTAG vs SWD	Same ARM debug access; SWD carries it on 2 pins instead of 4–5. Both reach breakpoints, memory and flash — pin count is the whole difference in practice.

PART IX • THE QUESTION BANK

Pure Q&A for the final hour — the answer, then the why, in one breath.

40 Objective Round — Answer + Why, One Breath Each

The short-answer questions that open most screenings. Give the answer, then the why, without pausing — the why is what gets scored.

Question	Answer — and the why
Can an ISR take parameters or return a value?	No — it is void f(void). Hardware calls it through the vector table; there is no caller to pass or receive anything. Exchange data via volatile flags or queues (§32).
Can you call printf or malloc inside an ISR?	Don't: both can block and are non-reentrant — an ISR landing mid-malloc corrupts the heap. Set a flag, defer the work to a task.
sizeof a pointer on a 32-bit MCU?	4 bytes regardless of the pointed-to type — a pointer stores an address, and addresses are word-sized.
sizeof('a')?	4 in C (character literals are int), 1 in C++. A pure language-lawyer check — knowing both marks real C time.
What does const volatile together mean?	Read-only to your code, yet changed by hardware behind your back — the exact type of a read-only status register (§5+§6 combined).
What is extern?	A declaration, not a definition: "storage lives in another file" — the linker resolves it. No memory is allocated at the extern line.
The four storage classes?	auto (block lifetime, the default), register (an obsolete hint), static (persistent + file-private — §39), extern (cross-file linkage).
0x12345678 stored little-endian at 0x2000 — byte order?	78 56 34 12 — LSB at the lowest address. Big-endian is the reverse; network order is big-endian, ARM is (usually) little.
Synchronous vs asynchronous communication?	Synchronous ships a clock wire (SPI, I2C); asynchronous agrees on timing in advance (UART) — which is why UART alone has baud-tolerance limits (§20).
Baud rate vs bit rate?	Baud = symbols per second, bit rate = bits per second. Equal only for two-level signaling like UART; multi-level links carry several bits per symbol.

Question	Answer — and the why
Simplex / half duplex / full duplex?	One-way only; two-way taking turns (RS-485); two-way simultaneously (SPI, a UART pair).
Interrupt vs polling — when each?	Interrupts for rare or urgent events — the CPU sleeps until needed. Polling for very frequent or jitter-critical sampling: its latency is boringly constant, which hard loops sometimes prefer.
Hard vs soft real time?	Hard: a missed deadline is a system failure (airbag). Soft: value merely degrades (a dropped video frame). The distinction is consequence, not speed.
Preemptive vs cooperative scheduling?	Preemptive: the kernel interrupts tasks by priority (RTOS default). Cooperative: tasks yield voluntarily — simpler, but one greedy loop starves everyone.
What is jitter? Drift?	Jitter is cycle-to-cycle variation of a period; drift is error accumulating over time. §30's timer-anchored loop kills both.
What is a critical section, on Cortex-M?	A region with interrupts masked (PRIMASK/BASEPRI) so a shared access completes atomically. Keep it instructions long, not loops — it adds directly to worst-case interrupt latency.
What is memory-mapped I/O?	Peripheral registers occupy addresses, so hardware is controlled by pointer writes — the reason volatile pointers to fixed addresses (§5) fill every HAL.
What is a HAL?	The hardware-abstraction layer: an API wall between intent (“toggle the pin”) and register-level how — the thing every portability answer leans on.
What is XIP?	Execute-in-place: run code straight from (Q)SPI flash without copying to RAM. Saves RAM; costs speed and determinism, so hot loops get relocated to RAM.
What is bit-banding?	A Cortex-M3/M4 alias region where every bit owns a word address — single-instruction atomic bit set/clear, sidestepping the read-modify-write race of §39's i++.
Why decoupling capacitors at every IC?	A local charge reservoir: supplies the nanosecond current spikes the supply trace cannot, and shunts high-frequency noise. 100 nF, as physically close to the pin as possible.
LED resistor: 3.3 V rail, $V_f = 2\text{ V}$, 10 mA?	$R = (V_{cc} - V_f) / I = 1.3\text{ V} / 10\text{ mA} = 130\ \Omega$. The formula is expected instantly — it is an arithmetic-under-pressure check.
Voltage divider output?	$V_{out} = V_{in} \cdot R_2 / (R_1 + R_2)$ — and volunteer the caveat: any load comparable to R_2 sags it, which is why dividers do not power things.
Setup and hold time?	Data must be stable for the setup window before a clock edge and the hold window after it. Violate either and the flip-flop may go metastable.
What is metastability — and the cure?	A flip-flop caught mid-decision outputs an undefined level for an unbounded moment. Standard cure: a 2-flip-flop synchronizer on every asynchronous input — buttons included (§39 debounce's cousin).
What is a BSP?	Board support package: the board-specific glue — pin muxing, clock tree, driver configuration — that lets generic firmware run on this particular PCB.
PWM duty cycle → average voltage?	$V_{avg} = D \cdot V_{cc}$ after filtering — which is why PWM plus an RC filter is the poor man's DAC, and why duty cycle IS the control variable in §33's loops.
Holonomic vs non-holonomic robot?	Holonomic can move in any direction instantly (omni/mecanum wheels); non-holonomic carries velocity constraints — a car cannot slide sideways. It changes which planner (§36) is even applicable.
What is odometry?	Dead-reckoning pose from wheel encoders: precise over meters, drifting with every slip — the exact reason §35's fusion exists.
How do you detect stack overflow?	Fill the stack with a pattern and check the high-water mark (RTOS built-ins), place a canary at the end, or — strongest — an MPU guard region that hard-faults on first touch (§7).

41 Subjective Round — Model Answers in Four Sentences

The open-ended questions where structure beats detail. Each model answer below is deliverable as-is — the parenthetical whys are what turn a list into an answer.

1. “A device in the field crashes randomly. Walk me through your debugging.”

First make the crash observable: a HardFault handler that saves PC, LR and the fault registers to noint RAM, plus RTOS stack high-water marks — because a random crash is almost always memory corruption, and corruption points backwards in time. Then narrow the suspects in probability order: stack overflow, an ISR/task race on shared data, a dangling pointer, and only then hardware (brownout, EMC). Bisect by disabling subsystems, and reproduce faster with stress (max interrupt rate, min stack sizes). The why behind the method: crashes are effects far from causes, so the job is instrumentation first, hypothesis second — never staring at code.

2. “How would you cut the power draw of a battery device?”

Architecture before tricks: make the system event-driven so the default state is deep sleep and everything wakes on interrupts — average current is dominated by how long you are awake, not how fast you compute (race-to-sleep). Then per-layer: gate unused peripheral clocks, drop the radio duty cycle (§25 — the radio usually owns the budget), let DMA move data while the core sleeps, and raise I2C pull-ups (§22). Finally, measure with a current profiler over a whole wake cycle — datasheet sleep numbers lie by omission (that one regulator or floating pin). The why: power is a time-integral, so the biggest wins are in the schedule, not the code.

3. “A brand-new board arrives from fab. How do you bring it up?”

Never flash first. Inspect, then power through a current-limited supply and verify every rail voltage — a solder bridge found at 50 mA is a component saved. Then establish the smallest possible sign of life: clock running, one LED blinking from bare register writes, then UART printf — because every later step needs a working output channel to complain through. Only then bring peripherals up one at a time, each proven in isolation (I2C address scan, SPI WHO_AM_I — §21). The why: bring-up is dependency ordering — each stage must be trustworthy before it becomes the tool that debugs the next.

4. “An I2C bus that worked yesterday is dead. Diagnose it.”

Scope before code: are both lines idling high? A line stuck low means a slave is holding SDA mid-byte — apply the 9-clock recovery ritual and a STOP (§22). If lines are high but no ACKs, run an address scan: a missing device points at power or address-pin strapping; all devices missing points at pull-ups or a shorted line. Check whether the master honors clock stretching — the classic works-then-dies cause with slow slaves. The why behind scope-first: I2C failures are electrical states, and code can only tell you what it *believes*, not what the wire is doing.

5. “How do you safely share data between an ISR and the main loop?”

Three tiers. A single aligned word: declare it volatile and rely on the bus’s atomic word access — sufficient and cheapest. Anything multi-word or read-modify-write: a critical section around the access, kept to a few instructions, or a lock-free single-producer ring buffer. Under an RTOS: hand the data through a queue or the §32 mailbox and let the kernel own the safety. Always name the why: volatile fixes *visibility* (compiler caching), never *atomicity* (§6) — interviewers are listening for exactly that distinction.

6. “How do you keep firmware portable across MCUs?”

Draw one line: application logic above, hardware below, and a thin HAL as the only crossing — the app calls `uart_send()`, never touches `USART1->DR`. Enforce it with fixed-width `stdint` types, wrapped compiler extensions, and ISRs that only pump events into the portable layer. If the product will outlive one chip, adopt an ecosystem built for this — Zephyr’s device tree moves the hardware description out of C entirely (§29). The why: porting cost is proportional to how many lines *know* which chip they run on; portability is minimizing that count from day one, because retrofitting the line is a rewrite.

7. “How do you test embedded software?”

Split by where the code can run. Pure logic (parsers, §33’s PID math, state machines) runs as unit tests on the host with mocked drivers — fast, in CI, on every commit. The hardware boundary gets integration tests on the real board: loopbacks, sensor sanity windows. The system level gets hardware-in-the-loop: simulate the plant, inject faults you could never safely produce physically (sensor dropout mid-motion), and regression-test timing. The why for HiL specifically: it makes the *rare and dangerous* cases repeatable — which is the entire hard part of embedded testing.

8. “Why is C still the embedded default — and when would you not use it?”

C survives because nothing is hidden: no allocator, no exceptions, no runtime doing work behind a line of code — so worst-case timing and memory are inspectable, which is the whole real-time contract. It maps one-to-one onto memory-mapped hardware and every vendor toolchain speaks it. I would deviate deliberately: a disciplined C++ subset

(no exceptions/RTTI, templates for zero-cost abstraction) on larger firmware, Rust where memory-safety bugs are the dominant risk and the toolchain is mature for the target, Python never on the MCU but always around it — tooling, tests, host simulation. The why interviewers want: the language choice is a determinism-and-visibility argument, not a fashion one.

9. “Design a safe over-the-air update.”

Three invariants, in order. Never erase the only working image: two slots (or app + recovery), the bootloader picks. Never run unverified bytes: CRC for corruption plus a signature for authenticity, checked *before* marking bootable. Never trust a new image’s word: it must confirm health after boot or a watchdog-backed rollback flag demotes it automatically (§12). Add resumable, power-fail-safe transfer state and you can yank power at any byte without a brick. The why in one word — anti-bricking: the design goal is that no single failure, including your own bad release, can strand the fleet.

10. “How would you localize a mobile robot?”

Layer fast-but-drifting sources under slow-but-absolute ones. Wheel odometry and gyro integration give smooth high-rate pose that drifts (§35); an EKF fuses them, weighting by uncertainty. Then anchor the drift with an absolute reference at lower rate — lidar scan matching against a map, visual fiducials, or UWB — whose corrections the filter blends in. State the failure modes unprompted: slip poisons odometry, glass and fog poison lidar — which is the argument for fusion itself: no single sensor’s lie is allowed to become the estimate. In ROS 2 terms: robot_localization’s EKF under, AMCL or SLAM above, all glued by the tf tree (§36).

11. “Choose the communication architecture for 12 sensors and 3 motor drivers on one robot.”

Segment by criticality, then by wiring reality. Motor drivers and E-stop: CAN — multi-drop, priority arbitration gives the stop command computable worst-case latency, and it shrugs off motor noise (§23). Slow ambient sensors: I2C, two wires and addressing for free, pull-ups sized for the speed (§22). Anything fast and point-to-point — an IMU at 1 kHz — SPI (§21). Uplink to the compute board: Ethernet/UDP or CAN-bridged. The why that scores: one bus for everything is the trap — each link is chosen by its worst case, not its average, and the E-stop’s worst case is the whole design driver.

12. “Tell me about the hardest bug you fixed.” — the structure, whatever your story is

Tell it in five beats, one sentence each: the **symptom** and its cost (“intermittent 2 s freeze in the field”); the **instrumentation** you added to see it (fault handler dump, GPIO-toggle timing, logs); the **hypothesis chain** — including one that was wrong and how the evidence killed it; the **root cause** named precisely (“priority inversion on the SPI mutex”, §39); and the **prevention** that outlived the fix (a lint rule, an assert, a test). The why: this question is never about the bug — it is an audit of your method, and the wrong-hypothesis beat is what proves the method is real.

How to use this in the room

Answer in three beats: **definition** → **why it exists** → **one trade-off or war story**. When you hear “always” or “never” in an option, it is almost certainly wrong. When stuck, reason out loud from first principles — interviewers grade the why-chain, not the recall.

The three answers that fail interviews

(1) “volatile makes it thread-safe” — it does not. (2) “malloc is fine, we have lots of RAM” — fragmentation, not size, is the killer. (3) “we’d just use a faster MCU” — real-time is about worst-case proof, not speed.